

Notes on Lambda Calculus

Spring 2017

The lambda calculus (or λ -calculus) was introduced by Alonzo Church and Stephen Cole Kleene in the 1930s to describe functions in an unambiguous and compact manner. Many real languages are directly based on the lambda calculus, such as Lisp, Scheme, Haskell, and ML. A key characteristic of these languages is that functions are values, just like integers and booleans are values: functions can be used as arguments to functions, and can be returned from functions. Those concepts have further made it into many other languages, including Python and JavaScript.

The name “lambda calculus” comes from the use of the Greek letter lambda (λ) in function definitions. (The letter lambda has no significance.) “Calculus” means a method of calculating by the symbolic manipulation of expressions.

Intuitively, a function is a rule for determining a value from an argument. Some examples of functions in mathematics are

$$f(x) = x^3$$

$$g(y) = y^3 - 2y^2 + 5y - 6.$$

(In the lambda notation introduced by Church and Kleene and used in the lambda calculus, these functions would be written $\lambda x. x^3$ and $\lambda y. y^3 - 2y^2 + 5y - 6$.)

1 Syntax

The pure λ -calculus contains just function definitions (called *abstractions*), variables, and function *application* (i.e., applying a function to an argument). If we add additional data types and operations (such as integers and addition), we have an *applied* λ -calculus. In the following text, we will sometimes assume that we have integers and addition in order to give more intuitive examples.

The syntax of the pure λ -calculus is defined as follows.

$e ::= x$	variable
$\lambda x. e$	abstraction
$e_1 e_2$	application

An abstraction $\lambda x. e$ is a function: variable x is the *argument*, and expression e is the *body* of the function. Note that the function $\lambda x. e$ doesn’t have a name. Assuming we have integers and arithmetic operations, the expression $\lambda y. y \times y$ is a function that takes an argument y and returns square of y .

An application $e_1 e_2$ requires that e_1 is (or evaluates to) a function, and then applies the function to the expression e_2 . For example, $(\lambda y. y \times y) 5$ is, intuitively, equal to 25, the result of applying the squaring function $\lambda y. y \times y$ to 5.

Here are some examples of lambda calculus expressions.

$\lambda x. x$	a lambda abstraction called the <i>identity function</i>
$\lambda x. (f (g x))$	another abstraction
$(\lambda x. x) 42$	an application
$\lambda y. \lambda x. x$	an abstraction that ignores its argument and returns the identity function

Lambda expressions extend as far to the right as possible. For example $\lambda x. x \lambda y. y$ is the same as $\lambda x. (x (\lambda y. y))$, and is not the same as $(\lambda x. x) (\lambda y. y)$. Application is left associative. For example $e_1 e_2 e_3$ is the same as $(e_1 e_2) e_3$. In general, use parentheses to make the parsing of a lambda expression clear if you are in doubt.

1.1 Variable binding and α -equivalence

An occurrence of a variable in an expression is either *bound* or *free*. An occurrence of a variable x in a term is bound if there is an enclosing $\lambda x. e$; otherwise, it is *free*. A *closed term* is one in which all identifiers are bound.

Consider the following term:

$$\lambda x. (x (\lambda y. y a) x) y$$

Both occurrences of x are bound, the first occurrence of y is bound, the a is free, and the last y is also free, since it is outside the scope of the λy .

If a program has some variables that are free, then you do not have a complete program as you do not know what to do with the free variables. Hence, a well formed program in lambda calculus is a closed term.

The symbol λ is a *binding operator*, as it binds a variable within some scope (i.e., some part of the expression): variable x is bound in e in the expression $\lambda x. e$.

The name of bound variables is not important. Consider the mathematical integrals $\int_0^7 x^2 dx$ and $\int_0^7 y^2 dy$. They describe the same integral, even though one uses variable x and the other uses variable y in their definition. The meaning of these integrals is the same: the bound variable is just a placeholder. In the same way, we can change the name of bound variables without changing the meaning of functions. Thus $\lambda x. x$ is the same function as $\lambda y. y$. Expressions e_1 and e_2 that differ only in the name of bound variables are called *α -equivalent* (“alpha equivalent”), sometimes written $e_1 =_\alpha e_2$.

1.2 Higher-order functions

In lambda calculus, functions are values: functions can take functions as arguments and return functions as results. In the pure lambda calculus, every value is a function, and every result is a function!

For example, the following function takes a function f as an argument, and applies it to the value 42.

$$\lambda f. f \ 42$$

This function takes an argument v and returns a function that applies its own argument (a function) to v .

$$\lambda v. \lambda f. (f \ v)$$

2 Semantics

2.1 β -equivalence

Application $(\lambda x. e_1) e_2$ applies the function $\lambda x. e_1$ to e_2 . In some ways, we would like to regard the expression $(\lambda x. e_1) e_2$ as equivalent to the expression e_1 where every (free) occurrence of x is replaced with e_2 . For example, we would like to regard $(\lambda y. y \times y) \ 5$ as equivalent to 5×5 .

We write $e_1\{e_2/x\}$ to mean expression e_1 with all free occurrences of x replaced with e_2 . There are several different notations to express this substitution, including $[x \mapsto e_2]e_1$ (used by Pierce), $[e_2/x]e_1$ (used by Mitchell), and $e_1[e_2/x]$ (used by Winskel).

Using our notation, we would like expressions $(\lambda x. e_1) e_2$ and $e_1\{e_2/x\}$ to be equivalent.

We call this equivalence, between $(\lambda x. e_1) e_2$ and $e_1\{e_2/x\}$, is called *β -equivalence*. Rewriting $(\lambda x. e_1) e_2$ into $e_1\{e_2/x\}$ is called a *β -reduction*. Given a lambda calculus expression, we may, in general, be able to perform β -reductions. This corresponds to executing a lambda calculus expression.

There may be more than one possible way to β -reduce an expression. Consider, for example, $(\lambda x. x + x) ((\lambda y. y) \ 5)$. We could use β -reduction to get either $((\lambda y. y) \ 5) + ((\lambda y. y) \ 5)$ or $(\lambda x. x + x) \ 5$. The order in which we perform β -reductions results in different semantics for the lambda calculus.

2.2 Evaluation strategies

There are many different evaluation strategies for the lambda calculus. The most permissive is full β -reduction, which allows any redex—i.e., any expression of the form $(\lambda x. e_1) e_2$ —to step to $e_1\{e_2/x\}$ at any time. It is defined formally by the following small-step operational semantics rules.

$$\frac{e_1 \rightarrow e'_1}{e_1 e_2 \rightarrow e'_1 e_2} \quad \frac{e_2 \rightarrow e'_2}{e_1 e_2 \rightarrow e_1 e'_2} \quad \frac{e \rightarrow e'}{\lambda x. e \rightarrow \lambda x. e'} \quad \beta\text{-REDUCTION} \frac{}{(\lambda x. e_1) e_2 \rightarrow e_1\{e_2/x\}}$$

A term e is said to be in *normal form* when it cannot be reduced any further, that is, when there is no e' such that $e \rightarrow e'$. It is convenient to say that term e *has* normal form e' if $e \rightarrow^* e'$ with e' in normal form.

Not every term has a normal form under full β -reduction. Consider the expression $(\lambda x. x x) (\lambda x. x x)$, which we will refer to as Ω for brevity. Let's try evaluating Ω .

$$\Omega = (\lambda x. x x) (\lambda x. x x) \rightarrow (\lambda x. x x) (\lambda x. x x) = \Omega$$


Evaluating Ω never reaches a term in normal form! It is an infinite loop!

When a term has a normal form, however, it never has more than one. This is not a given, because clearly the full β -reduction strategy is non-deterministic. Look at the term $(\lambda x. \lambda y. y) \Omega (\lambda z. z)$, for example. It has two redexes in it, the one with abstraction λx , and the one inside Ω . But this nondeterminism is well behaved: when different sequences of reduction reach a normal form (they need not) those normal forms are equal.

Formally, full β -reduction is confluent in the following sense:

Theorem 1 (Confluence). *If $e \rightarrow^* e_1$ and $e \rightarrow^* e_2$ then there exists e' such that $e_1 \rightarrow^* e'$ and $e_2 \rightarrow^* e'$.*

Confluence can be depicted graphically as follows (where \rightarrow is used to represent \rightarrow^*):



Confluence is often also called the Church-Rosser property. It is not an easy result to prove. (It would make sense for it to be a proof by induction on the multi-step reduction \rightarrow^* . Try it, and see where you get stuck.)

Corollary 1. *If $e \rightarrow^* e_1$ and $e \rightarrow^* e_2$ and both e_1 and e_2 are in normal form, then $e_1 = e_2$.*

Proof. An easy consequence of confluence. □

Other evaluation strategies are possible, which impose a deterministic order on the reductions. For example, *normal order evaluation* uses the full β -reduction rules, except imposes the order that the left-most redex—that is, the redex in which the leading λ appears left-most in the term—is always reduced first. Normal order evaluation guarantees that if a term has a normal form, applying reductions in normal order will eventually yield that normal form.

Normal order evaluation allows reducing redexes inside abstractions, which may strike you as odd if you rely on your programmer's intuition: a function definition does not simply reduce its body, unprompted. That's because most programming languages use reduction strategies that when put in lambda calculus terms do not perform reductions inside abstractions.

Two common evaluations strategies that occur in programming languages are call-by-value and call-by-name.

Call-by-value (or CBV) evaluation strategy is more restrictive: it only allows an application to reduce after its argument has been reduced to a value and does not allow evaluation under a λ . That is, given an application $(\lambda x. e_1) e_2$, CBV semantics makes sure that e_2 is a value before calling the function.

So, what is a value? In the pure lambda calculus, any abstraction is a value. Remember, an abstraction $\lambda x. e$ is a function; in the pure lambda calculus, the only values are functions. In an *applied lambda calculus* with integers and arithmetic operations, values also include integers. Intuitively, a value is an expression that can not be reduced/executed/simplified any further.

We can give small-step operational semantics for call-by-value execution of the lambda calculus. Here, v can be instantiated with any value (e.g., a function).

$$\frac{e_1 \longrightarrow e'_1}{e_1 e_2 \longrightarrow e'_1 e_2} \qquad \frac{e \longrightarrow e'}{v e \longrightarrow v e'} \qquad \beta\text{-REDUCTION} \frac{}{(\lambda x. e) v \longrightarrow e\{v/x\}}$$

We can see from these rules that, given an application $e_1 e_2$, we first evaluate e_1 until it is a value, then we evaluate e_2 until it is a value, and then we apply the function to the value—a β -reduction.

Let's consider some examples. (These examples use an applied lambda calculus that also includes reduction rules for arithmetic expressions.)

$$\begin{aligned} (\lambda x. \lambda y. y x) (5 + 2) \lambda x. x + 1 &\longrightarrow (\lambda x. \lambda y. y x) 7 \lambda x. x + 1 \\ &\longrightarrow (\lambda y. y 7) \lambda x. x + 1 \\ &\longrightarrow (\lambda x. x + 1) 7 \\ &\longrightarrow 7 + 1 \\ &\longrightarrow 8 \end{aligned}$$

$$\begin{aligned} (\lambda f. f 7) ((\lambda x. x x) \lambda y. y) &\longrightarrow (\lambda f. f 7) ((\lambda y. y) (\lambda y. y)) \\ &\longrightarrow (\lambda f. f 7) (\lambda y. y) \\ &\longrightarrow (\lambda y. y) 7 \\ &\longrightarrow 7 \end{aligned}$$

Call-by-name (or CBN) semantics are more permissive than CBV, but less permissive than full β -reduction. CBN semantics applies the function as soon as possible. The small-step operational semantics are a little simpler, as they do not need to ensure that the expression to which a function is applied is a value.

$$\frac{e_1 \longrightarrow e'_1}{e_1 e_2 \longrightarrow e'_1 e_2} \qquad \beta\text{-REDUCTION} \frac{}{(\lambda x. e_1) e_2 \longrightarrow e_1\{e_2/x\}}$$

Let's consider the same examples we used for CBV.

$$\begin{aligned} (\lambda x. \lambda y. y x) (5 + 2) \lambda x. x + 1 &\longrightarrow (\lambda y. y (5 + 2)) \lambda x. x + 1 \\ &\longrightarrow (\lambda x. x + 1) (5 + 2) \\ &\longrightarrow (5 + 2) + 1 \\ &\longrightarrow 7 + 1 \\ &\longrightarrow 8 \end{aligned}$$

$$\begin{aligned} (\lambda f. f 7) ((\lambda x. x x) \lambda y. y) &\longrightarrow ((\lambda x. x x) \lambda y. y) 7 \\ &\longrightarrow ((\lambda y. y) (\lambda y. y)) 7 \\ &\longrightarrow (\lambda y. y) 7 \\ &\longrightarrow 7 \end{aligned}$$

Note that the answers are the same, but the order of evaluation is different. (Later we will see languages where the order of evaluation is important, and may result in different answers.)

One way in which CBV and CBN differ is when arguments to functions have no normal forms. For instance, consider the following term:

$$(\lambda x.(\lambda y.y)) \Omega$$

If we use CBV semantics to evaluate the term, we must reduce Ω to a value before we can apply the function. But Ω never evaluates to a value, so we can never apply the function. Under CBV semantics, this term does not have a normal form.

If we use CBN semantics, then we can apply the function immediately, without needing to reduce the actual argument to a value. We have

$$(\lambda x.(\lambda y.y)) \Omega \longrightarrow_{\text{CBN}} \lambda y.y$$

CBV and CBN are common evaluation orders; many programming languages use CBV semantics. So-called “lazy” languages, such as Haskell, typically use Call-by-need semantics, a more efficient semantics similar to CBN in that it does not evaluate actual arguments unless necessary. However, Call-by-need semantics ensures that arguments are evaluated at most once.

3 Lambda calculus encodings

The pure lambda calculus contains only functions as values. It is not exactly easy to write large or interesting programs in the pure lambda calculus. We can however encode objects, such as booleans, and integers.

3.1 Booleans

We want to encode constants and operators for booleans. That is, we want to define functions *TRUE*, *FALSE*, *AND*, *IF*, and other operators such that the expected behavior holds, for example:

$$\begin{aligned} \text{AND } \text{TRUE } \text{FALSE} &= \text{FALSE} \\ \text{IF } \text{TRUE } e_1 \ e_2 &= e_1 \\ \text{IF } \text{FALSE } e_1 \ e_2 &= e_2 \end{aligned}$$

Let’s start by defining *TRUE* and *FALSE* as follows.

$$\begin{aligned} \text{TRUE} &\triangleq \lambda x. \lambda y. x \\ \text{FALSE} &\triangleq \lambda x. \lambda y. y \end{aligned}$$

Thus, both *TRUE* and *FALSE* take two arguments, *TRUE* returns the first, and *FALSE* returns the second.

The function *IF* should behave like $\lambda b. \lambda t. \lambda f. \text{if } b = \text{TRUE} \text{ then } t \text{ else } f$. The definitions for *TRUE* and *FALSE* make this very easy.

$$\text{IF} \triangleq \lambda b. \lambda t. \lambda f. b \ t \ f$$

Definitions of other operators are also straightforward.

$$\begin{aligned} \text{NOT} &\triangleq \lambda b. b \ \text{FALSE} \ \text{TRUE} \\ \text{AND} &\triangleq \lambda b_1. \lambda b_2. b_1 \ b_2 \ \text{FALSE} \\ \text{OR} &\triangleq \lambda b_1. \lambda b_2. b_1 \ \text{TRUE} \ b_2 \end{aligned}$$

3.2 Church numerals

Church numerals encode the natural number n as a function that takes f and x , and applies f to x n times.

$$\begin{aligned}\bar{0} &\triangleq \lambda f. \lambda x. x \\ \bar{1} &= \lambda f. \lambda x. f \ x \\ \bar{2} &= \lambda f. \lambda x. f \ (f \ x) \\ \text{SUCC} &\triangleq \lambda n. \lambda f. \lambda x. f \ (n \ f \ x)\end{aligned}$$

In the definition for *SUCC*, the expression $n \ f \ x$ applies f to x n times (assuming that variable n is the Church encoding of the natural number n). We then apply f to the result, meaning that we apply f to x $n + 1$ times.

Given the definition of *SUCC*, we can easily define addition. Intuitively, the natural number $n_1 + n_2$ is the result of apply the successor function n_1 times to n_2 .

$$\text{ADD} \triangleq \lambda n_1. \lambda n_2. n_1 \ \text{SUCC} \ n_2$$

Similarly, we can define multiplication, by noting $n_1 \times n_2$ is the result of applying n_1 times to 0 that function that adds n_2 to its input. The latter can be obtained by considering *ADD* n_2 , and thus

$$\text{MUL} \triangleq \lambda n_1. \lambda n_2. n_1 \ (\text{ADD} \ n_2) \ \bar{0}$$

It is a lot more challenging to define subtraction. The difficulty is defining a function that takes a Church numeral representing n and returning its predecessor, the Church numeral representing $n - 1$. It is possible to define such a function *PRED*. It is a difficult exercise, but the answer is easy enough to find online. Here's an intuition to get you started: think about enumerating the pairs $(0, 1), (1, 2), (2, 3), (3, 4), \dots$. Finding the predecessor of n amounts to finding the n th pair in this enumeration, and looking at the first component of the pair.

How do we encode pairs, though? Funny you should ask...

3.3 Pairs

A pair (a, b) is a packaging up of two elements a and b in such a way that you can treat the package as a single unit, and retrieve both a and b later. This means that we're looking for a functions *PAIR*, *FIRST*, and *SECOND* with the property that:

$$\begin{aligned}\text{FIRST}(\text{PAIR} \ a \ b) &= a \\ \text{SECOND}(\text{PAIR} \ a \ b) &= b\end{aligned}$$

There are several encodings possible that satisfy this specification. The easiest is probably to encode a pair as a function that expects a “selector” and applies it to both elements of the pair. *FIRST* and *SECOND* then simply supply the appropriate selector to the pair.

$$\begin{aligned}\text{PAIR} &= \lambda a. \lambda b. \lambda s. s \ a \ b \\ \text{FIRST} &= \lambda p. p \ (\lambda x. \lambda y. x) \\ \text{SECOND} &= \lambda p. p \ (\lambda x. \lambda y. y)\end{aligned}$$

It is easy to generalize this encoding to arbitrary k -tuples, and even to arbitrary-sized lists with constructors *NIL* and *CONS* and accessors *HEAD* and *TAIL*.

4 Recursion and the fixed-point combinators

We can write nonterminating functions, as we saw with the expression Ω . We can also write recursive functions that terminate. However, one complication is how we express this recursion.

Let's consider how we would like to define a function that computes factorials.

$$FACT \triangleq \lambda n. IF (ISZERO\ n) 1 (MUL\ n\ (FACT\ (PRED\ n)))$$

(We have not defined the predicate *ISZERO*. It is an easy exercise though.)

In slightly more readable notation (and we will see next lecture how we can translate more readable notation into appropriate expressions):

$$FACT \triangleq \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times FACT\ (n - 1)$$

Here, like in the definitions we gave above, the name *FACT* is simply meant to be shorthand for the expression on the right-hand side of the equation. But *FACT* appears on the right-hand side of the equation as well! This is not a definition, it's a recursive equation.

4.1 Recursion removal trick

We can perform a “trick” to define a function *FACT* that satisfies the recursive equation above. First, let's define a new function *FACT'* that looks like *FACT*, but takes an additional argument *f*. We assume that the function *f* will be instantiated with an actual parameter of... *FACT'*.

$$FACT' \triangleq \lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f\ f\ (n - 1))$$

Note that when we call *f*, we pass it a copy of itself, preserving the assumption that the actual argument for *f* will be *FACT'*.

Now we can define the factorial function *FACT* in terms of *FACT'*.

$$FACT \triangleq FACT' FACT'$$

Let's try evaluating *FACT* applied to an integer.

$FACT\ 3 = (FACT'\ FACT')\ 3$	Definition of <i>FACT</i>
$= ((\lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f\ f\ (n - 1)))\ FACT')\ 3$	Definition of <i>FACT'</i>
$\longrightarrow (\lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (FACT'\ FACT'\ (n - 1)))\ 3$	Application to <i>FACT'</i>
$\longrightarrow \text{if } 3 = 0 \text{ then } 1 \text{ else } 3 \times (FACT'\ FACT'\ (3 - 1))$	Application to <i>n</i>
$\longrightarrow 3 \times (FACT'\ FACT'\ (3 - 1))$	Evaluating if
$\longrightarrow \dots$	
$\longrightarrow 3 \times 2 \times 1 \times 1$	
$\longrightarrow^* 6$	

So we now have a technique for writing a recursive function *f*: write a function *f'* that explicitly takes a copy of itself as an argument, and then define $f \triangleq f'\ f'$.

4.2 Fixed point combinators

There is another way of writing recursive functions: expressing the recursive function as the fixed point of some other, higher-order function, and then finding that fixed point.

Let's consider the factorial function again. The factorial function *FACT* is a fixed point of the following function.

$$G \triangleq \lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f\ (n - 1))$$

(Recall that if *g* is a fixed point of *G*, then we have $G\ g = g$.)

So if we had some way of finding a fixed point of G , we would have a way of defining the factorial function $FACT$.

There are such “fixed point operators,” and the (infamous) Y combinator is one of them. Thus, we can define the factorial function $FACT$ to be simply $Y\ G$, the fixed point of G .

(A *combinator* is simply a closed lambda term; it is a higher-order function that uses only function application and other combinators to define a result from its arguments; our functions $SUCC$ and ADD are examples of combinators. It is possible to define programs using only combinators, thus avoiding the use of variables completely.)

The Y combinator is defined as

$$Y \triangleq \lambda f. (\lambda x. f\ (x\ x))\ (\lambda x. f\ (x\ x)).$$

It was discovered by Haskell Curry, and is one of the simplest fixed-point combinators.

The fixed point of the higher-order function G is equal to $G\ (G\ (G\ (G\ (G\ \dots))))$. Intuitively, the Y combinator unrolls this equality, as needed. Let’s see it in action, on our function G , where

$$G = \lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f\ (n - 1))$$

and the factorial function is the fixed point of G . (We will use CBN semantics; see the note below.)

$$\begin{aligned} FACT &= Y\ G \\ &= (\lambda f. (\lambda x. f\ (x\ x))\ (\lambda x. f\ (x\ x)))\ G && \text{Definition of } Y \\ &\longrightarrow (\lambda x. G\ (x\ x))\ (\lambda x. G\ (x\ x)) \\ &\longrightarrow G\ (\lambda x. G\ (x\ x))\ (\lambda x. G\ (x\ x)) \end{aligned}$$

Here, note that $(\lambda x. G\ (x\ x))\ (\lambda x. G\ (x\ x))$ was the result of beta-reducing $Y\ G$. That is $(\lambda x. G\ (x\ x))\ (\lambda x. G\ (x\ x))$ is β -equivalent to $Y\ G$ which is equal to $FACT$. So we will rewrite the expression as follows.

$$\begin{aligned} &=_{\beta} G\ (FACT) \\ &= (\lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f\ (n - 1)))\ FACT && \text{Definition of } G \\ &\longrightarrow \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (FACT\ (n - 1)) \end{aligned}$$

Note that the Y combinator works under CBN semantics, but not CBV. What happens when we evaluate $Y\ G$ under CBV? Have a try and see. There is a variant of the Y combinator, Z , that works under CBV semantics. It is defined as

$$Z \triangleq \lambda f. (\lambda x. f\ (\lambda y. x\ x\ y))\ (\lambda x. f\ (\lambda y. x\ x\ y)).$$

There are many (indeed infinite) fixed-point combinators. To gain some more intuition for fixed-point combinators, let’s derive the Turing fixed-point combinator, discovered by Alan Turing, and denoted by Θ .

Suppose we have a higher-order function f , and want the fixed point of f . We know that $\Theta\ f$ is a fixed point of f , so we have

$$\Theta\ f = f\ (\Theta\ f).$$

This means, that we can write the following recursive equation for Θ .

$$\Theta = \lambda f. f\ (\Theta\ f)$$

Now we can use the recursion removal trick we described earlier! Let’s define $\Theta' = \lambda t. \lambda f. f\ (t\ t\ f)$, and define

$$\begin{aligned} \Theta &\triangleq \Theta'\ \Theta' \\ &= (\lambda t. \lambda f. f\ (t\ t\ f))\ (\lambda t. \lambda f. f\ (t\ t\ f)) \end{aligned}$$

Let's try out the Turing combinator on our higher-order function G that we used to define $FACT$. Again, we will use CBN semantics.

$$\begin{aligned}
FACT &= \Theta \ G \\
&= ((\lambda t. \lambda f. f \ (t \ t \ f)) \ (\lambda t. \lambda f. f \ (t \ t \ f))) \ G \\
&\rightarrow (\lambda f. f \ ((\lambda t. \lambda f. f \ (t \ t \ f)) \ (\lambda t. \lambda f. f \ (t \ t \ f)) \ f)) \ G \\
&\rightarrow G \ ((\lambda t. \lambda f. f \ (t \ t \ f)) \ (\lambda t. \lambda f. f \ (t \ t \ f)) \ G) \\
&= G \ (\Theta \ G) && \text{for brevity} \\
&= (\lambda f. \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f \ (n - 1))) \ (\Theta \ G) && \text{Definition of } G \\
&\rightarrow \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times ((\Theta \ G) \ (n - 1)) \\
&= \lambda n. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (FACT \ (n - 1))
\end{aligned}$$

5 Definitional translation

We saw that the denotational semantics of IMP defined the meaning of IMP commands as mathematical functions from stores to stores. We can describe denotational semantics as a form of compilation, from IMP to mathematics. We now consider definitional translation, where we define the meaning of language constructs by translation to another language. This is a form of denotational semantics, but instead of the target language being mathematics, it is a simpler programming language. Note that definitional translation does not necessarily produce clean or efficient code; rather, it defines the meaning of the source language in terms of the target language.

We now consider a number of language features we can add to the lambda calculus, define an operational semantics for them, and then give an alternate semantics by translation to the simpler language that is the lambda calculus without additional features. We first introduce *evaluation contexts* to help us present the new language features succinctly.

5.1 Evaluation contexts

Recall the syntax and CBV operational semantics for the lambda calculus.

$$\begin{aligned}
e &::= x \mid \lambda x. e \mid e_1 \ e_2 \\
\frac{e_1 \rightarrow e'_1}{e_1 \ e_2 \rightarrow e'_1 \ e_2} & \quad \frac{e \rightarrow e'}{v \ e \rightarrow v \ e'} && \beta\text{-REDUCTION} \frac{}{(\lambda x. e) \ v \rightarrow e\{v/x\}}
\end{aligned}$$

Of the operational semantics rules, only the β -reduction rule told us how to “reduce” an expression; the other two rules were simply telling us the order to evaluate expressions in, i.e., first evaluate the left hand side of an application to a value, then evaluate the right hand side of an application to a value. The operational semantics of many of the languages we will consider have this feature: there are two kinds of rules, one kind specifying evaluation order, and the other kind specifying the “interesting” reductions.

Evaluation contexts provide us with a mechanism to separate out these two kinds of rules. An evaluation context E (sometimes written $E[\cdot]$) is an expression with a “hole” in it, that is with a single occurrence of the special symbol $[\cdot]$ (called the “hole”) in place of a subexpression. Evaluation contexts are defined using a BNF grammar that is similar to the grammar used to define the language. The following grammar defines evaluation contexts for the pure CBV lambda calculus.

$$E ::= [\cdot] \mid E \ e \mid v \ E$$

We write $E[e]$ to mean the evaluation context E where the hole has been replaced with the expression e . The following are examples of evaluation contexts, and evaluation contexts with the hole filled in by an

expression.

$$\begin{aligned}
E_1 &= [\cdot] (\lambda x. x) & E_1[\lambda y. y \ y] &= (\lambda y. y \ y) \ \lambda x. x \\
E_2 &= (\lambda z. z \ z) [\cdot] & E_2[\lambda x. \lambda y. x] &= (\lambda z. z \ z) (\lambda x. \lambda y. x) \\
E_3 &= ([\cdot] \ \lambda x. x \ x) ((\lambda y. y) (\lambda y. y)) & E_3[\lambda f. \lambda g. f \ g] &= ((\lambda f. \lambda g. f \ g) \ \lambda x. x \ x) ((\lambda y. y) (\lambda y. y))
\end{aligned}$$

Using evaluation contexts, we can define the evaluation semantics for the pure CBV lambda calculus with just two rules, one for evaluation contexts, and one for β -reduction.

$$\frac{e \longrightarrow e'}{E[e] \longrightarrow E[e']} \quad \beta\text{-REDUCTION} \frac{}{(\lambda x. e) \ v \longrightarrow e\{v/x\}}$$

Note that the evaluation contexts for the CBV lambda calculus ensure that we evaluate the left hand side of an application to a value, and then evaluate the right hand side of an application to a value before applying β -reduction.

We can specify the operational semantics of CBN lambda calculus using evaluation contexts:

$$E ::= [\cdot] \mid E \ e \quad \frac{e \longrightarrow e'}{E[e] \longrightarrow E[e']} \quad \beta\text{-REDUCTION} \frac{}{(\lambda x. e_1) \ e_2 \longrightarrow e_1\{e_2/x\}}$$

We'll see the benefit of evaluation contexts as we see languages with more syntactic constructs.

5.2 Multi-argument functions and currying

Our syntax for functions restricted us to function that have a single argument: $\lambda x. e$. We could define a language that allows functions to have multiple arguments.

$$e ::= x \mid \lambda \langle x_1, \dots, x_n \rangle. e \mid e_0 \ \langle e_1, \dots, e_n \rangle$$

Here, a function $\lambda \langle x_1, \dots, x_n \rangle. e$ takes n arguments, with names x_1 through x_n . In a multi-argument application $e_0 \ \langle e_1, \dots, e_n \rangle$, we expect e_0 to evaluate to an n -argument function, and e_1, \dots, e_n are the arguments that will give the function.

We can define a CBV operational semantics for the multi-argument lambda calculus as follows.

$$E ::= [\cdot] \mid E \ \langle e_1, \dots, e_n \rangle \mid v_0 \ \langle v_1, \dots, v_{i-1}, E, e_{i+1}, \dots, e_n \rangle$$

$$\frac{e \longrightarrow e'}{E[e] \longrightarrow E[e']} \quad \beta\text{-REDUCTION} \frac{}{(\lambda \langle x_1, \dots, x_n \rangle. e_0) \ \langle v_1, \dots, v_n \rangle \longrightarrow e_0\{v_1/x_1\}\{v_2/x_2\} \dots \{v_n/x_n\}}$$

The evaluation contexts ensure that we evaluate a multi-argument application $e_0 \ \langle e_1, \dots, e_n \rangle$ by evaluating each expression from left to right down to a value.

Now, the multi-argument lambda calculus isn't any more expressive than the pure lambda calculus. We can show this by showing how any multi-argument lambda calculus program can be translated into an equivalent pure lambda calculus program. We define a translation function $\mathcal{T}[\cdot]$ that takes an expression in the multi-argument lambda calculus and returns an equivalent expression in the pure lambda calculus. That is, if e is a multi-argument lambda calculus expression, $\mathcal{T}[e]$ is a pure lambda calculus expression.

We define the translation as follows.

$$\begin{aligned}
\mathcal{T}[x] &= x \\
\mathcal{T}[\lambda \langle x_1, \dots, x_n \rangle. e] &= \lambda x_1. \dots \lambda x_n. \mathcal{T}[e] \\
\mathcal{T}[e_0 \ \langle e_1, \dots, e_n \rangle] &= (\dots ((\mathcal{T}[e_0] \ \mathcal{T}[e_1]) \ \mathcal{T}[e_2]) \dots \mathcal{T}[e_n])
\end{aligned}$$

This process of rewriting a function that takes multiple arguments as a chain of functions that each take a single argument is called *currying*. Consider a mathematical function that takes two arguments, the first from domain A and the second from domain B , and returns a result from domain C . We could describe this function, using mathematical notation for domains of functions, as being an element of $A \times B \rightarrow C$. Currying this function produces a function that is an element of $A \rightarrow (B \rightarrow C)$. That is, the curried version of the function takes an argument from domain A , and returns a function that takes an argument from domain B and produces a result of domain C .

5.3 Products and let

We introduce two useful language features to the lambda calculus: products and let expressions.

A product is a pair of expressions (e_1, e_2) . If e_1 and e_2 are both values, then we regard the product as also being a value. (That is, we cannot further evaluate a product if both elements are values.)

Given, a product, we can access the first or second element using the operators $\#1$ and $\#2$ respectively. That is, $\#1 (v_1, v_2) \rightarrow v_1$ and $\#2 (v_1, v_2) \rightarrow v_2$. (Other common notation for projection includes π_1 and π_2 , and **fst** and **snd**.)

More formally, we define the syntax of lambda calculus with products and let expressions as follows. Values in this language are either functions or pairs of values.

$$\begin{aligned} e &::= x \mid \lambda x. e \mid e_1 e_2 \mid (e_1, e_2) \mid \#1 e \mid \#2 e \mid \text{let } x = e_1 \text{ in } e_2 \\ v &::= \lambda x. e \mid (v_1, v_2) \end{aligned}$$

We define a small-step CBV operational semantics for the language using evaluation contexts.

$$\begin{aligned} E &::= [\cdot] \mid E e \mid v E \mid (E, e) \mid (v, E) \mid \#1 E \mid \#2 E \mid \text{let } x = E \text{ in } e_2 \\ \frac{e \rightarrow e'}{E[e] \rightarrow E[e']} & \qquad \beta\text{-REDUCTION} \frac{}{(\lambda x. e) v \rightarrow e\{v/x\}} \\ \frac{}{\#1 (v_1, v_2) \rightarrow v_1} & \qquad \frac{}{\#2 (v_1, v_2) \rightarrow v_2} \\ \frac{}{\text{let } x = v \text{ in } e \rightarrow e\{v/x\}} & \end{aligned}$$

We can give an equivalent semantics by translation to the pure CBV lambda calculus. We encode a pair (e_1, e_2) using the same encoding we saw earlier in lecture, as a value that takes a function f , and applies f to v_1 and v_2 , where v_1 and v_2 are the result of evaluating e_1 and e_2 respectively. The projection operators pass a function to the encoding of pairs that selects either the first or second element as appropriate.

Note also that the expression $\text{let } x = e_1 \text{ in } e_2$ is equivalent to the application $(\lambda x. e_2) e_1$.

$$\begin{aligned} \mathcal{T}[x] &= x \\ \mathcal{T}[\lambda x. e] &= \lambda x. \mathcal{T}[e] \\ \mathcal{T}[e_1 e_2] &= \mathcal{T}[e_1] \mathcal{T}[e_2] \\ \mathcal{T}[(e_1, e_2)] &= (\lambda x. \lambda y. \lambda f. f x y) \mathcal{T}[e_1] \mathcal{T}[e_2] \\ \mathcal{T}[\#1 e] &= \mathcal{T}[e] (\lambda x. \lambda y. x) \\ \mathcal{T}[\#2 e] &= \mathcal{T}[e] (\lambda x. \lambda y. y) \\ \mathcal{T}[\text{let } x = e_1 \text{ in } e_2] &= (\lambda x. \mathcal{T}[e_2]) \mathcal{T}[e_1] \end{aligned}$$

5.4 CBN to CBV

We've seen semantics for both the call-by-name lambda calculus and the call-by-value lambda calculus. We can translate a call-by-name program into a call-by-value program. In CBV, arguments to functions are evaluated before the function is applied; in CBN, functions are applied as soon as possible. In the translation, we delay the evaluation of arguments by wrapping them in a function. This is called a *thunk*: wrapping a computation in a function to delay its evaluation.

Since arguments to functions are turned into thunks, when we want to use an argument in a function body, we need to evaluate the thunk. We do so by applying the thunk (which is simply a function); it doesn't matter what we apply the thunk to, since the thunk's argument is never used.

$$\begin{aligned}\mathcal{T}[\![x]\!] &= x \ (\lambda y. y) \\ \mathcal{T}[\![\lambda x. e]\!] &= \lambda x. \mathcal{T}[\![e]\!] \\ \mathcal{T}[\![e_1 \ e_2]\!] &= \mathcal{T}[\![e_1]\!] \ (\lambda z. \mathcal{T}[\![e_2]\!]) \quad z \text{ is not a free variable of } e_2\end{aligned}$$

5.5 Adequacy of translation

We've presented several translations of languages. In each case, we had a semantics defined for both the source and target language. We would like the translation to be correct, that is, to preserve the meaning of source programs.

More precisely, we would like an expression e in the source language to evaluate to a value v if and only if the translation of e evaluates to a value v' such that v' is “equal to” v .

What exactly it means for v' to be “equal to” v will depend on the translation. Sometimes, it will mean that v' is the translation of v ; other times, it will mean that v' is somehow equivalent to the translation of v . In particular, we often need to define equivalence on functions. One possible solution is that two functions are equivalent if they agree on the result when applied to any value of a base type (e.g., integers or booleans). The idea is that if two functions disagree when passed a more complex value (say, a function), then we could write a program that uses these functions to produce functions that disagree on values of base types.

There are two criteria for a translation to be *adequate*: soundness and completeness. For clarity, let's suppose that $\mathbf{Exp}_{\text{src}}$ is the set of source language expressions, and that $\longrightarrow_{\text{src}}$ and $\longrightarrow_{\text{trg}}$ are the evaluation relations for the source and target languages respectively.

A translation is sound if every target evaluation represents a source evaluation:

$$\text{Soundness: } \forall e \in \mathbf{Exp}_{\text{src}}. \text{ if } \mathcal{T}[\![e]\!] \longrightarrow_{\text{trg}}^* v' \text{ then } \exists v. e \longrightarrow_{\text{src}}^* v \text{ and } v' \text{ equivalent to } v$$

A translation is complete if every source evaluation has a target evaluation.

$$\text{Completeness: } \forall e \in \mathbf{Exp}_{\text{src}}. \text{ if } e \longrightarrow_{\text{src}}^* v \text{ then } \exists v'. \mathcal{T}[\![e]\!] \longrightarrow_{\text{trg}}^* v' \text{ and } v' \text{ equivalent to } v$$

6 Continuations

So far we have seen a number of language features that extend lambda calculus, and have translated many of these into the pure lambda calculus, using a direct semantic translation. That is, the control structure of the source language translated into the corresponding control structure in the target language:

$$\begin{aligned}\mathcal{T}[\![\lambda x. e]\!] &= \lambda x. \mathcal{T}[\![e]\!] \\ \mathcal{T}[\![e_1 \ e_2]\!] &= \mathcal{T}[\![e_1]\!] \ \mathcal{T}[\![e_2]\!]\end{aligned}$$

This style of translation works well when the source language is similar to the target language. However, when the control structures of the source and target languages differ considerably, it doesn't work as well.

In this lecture, we'll learn about *continuations* and *continuation-passing style*. These can be used to define the semantics of control-flow constructs such as exceptions.

Continuations are a programming technique that may be used directly by a programmer, or used in program transformations by a compiler.

Intuitively, a continuation represents “the rest of the program.” Consider the program

if `foo < 10` then `32 + 6` else `7 + bar`

and consider the evaluation of the expression `foo < 10`. When we finish evaluating this subexpression, we will evaluate the if statement, and then evaluate the appropriate branch. The *continuation* of the subexpression `foo < 10` is the rest of the computation that will occur after we evaluate the subexpression. We can write this continuation as a function that takes the result of the subexpression:

$(\lambda y. \text{if } y \text{ then } 32 + 6 \text{ else } 7 + \text{bar}) (\text{foo} < 10)$

The evaluation order and result of this program will be the same as the original expression; the difference is that we extracted the continuation of the subexpression in to a function.

The nice thing about continuations is that it makes the control explicit, and this is especially useful in the case of functional programs, where control is not explicit otherwise. In fact, we can rewrite a program to make continuations more explicit. The idea is to turn evaluation contexts (an ingredient of the semantic description) into continuations (a concrete artifact within a program). Roughly, we will turn an evaluation context $E[\cdot]$ into a function k_E that expects a value to put into the hole, so that $E[e]$ can be interpreted as calling the continuation as $k_E e$.

Here’s an example. Let’s convert the following applied lambda calculus program

$((1 + 2) + 3) + 4$

to use explicit continuations, by repeatedly “peeling away” the evaluation contexts.

The evaluation context used to isolate the first subexpression to evaluate is

$E_1 = ([\cdot] + 3) + 4$

If k_1 is the continuation corresponding to E_1 , then we can write

$$\begin{aligned} e &= E_1[1 + 2] \\ &= k_1 (1 + 2) \end{aligned}$$

where $k_1 = \lambda a. E_1[a]$. Since $E_1[a] = E_2[a + 3]$ where $E_2 = [\cdot] + 4$, we can take:

$$\begin{aligned} k_1 &= \lambda a. E_1[a] \\ &= \lambda a. E_2[a + 3] \\ &= \lambda a. k_2 (a + 3) \end{aligned}$$

where $k_2 = \lambda b. E_2[b]$. Since $E_2[b] = E_3[b + 4]$ where $E_3 = [\cdot]$ we can take

$$\begin{aligned} k_2 &= \lambda b. E_2[b] \\ &= \lambda b. E_3[b + 4] \\ &= \lambda b. k_3 (b + 4) \end{aligned}$$

where

$$\begin{aligned} k_3 &= \lambda c. E_3[c] \\ &= \lambda c. c \end{aligned}$$

Expanding everything out, we have

$e = (\lambda a. (\lambda b. (\lambda c. c) (b + 4)) (a + 3)) (1 + 2)$

This is nigh unreadable. But if we recall that $\text{let } x = e \text{ in } e'$ is just syntactic sugar for $(\lambda x. e') e$, we can actually rewrite the above as

```
let a = 1 + 2 in
let b = a + 3 in
let c = b + 4 in
c
```

A key feature of this kind of transformation is now clear from the resulting code above: after the transformation, primitive operations are only applied to values or to variables. The evaluation order (and hence the control flow) is made fully explicit by the sequence of applications.

As an interesting aside, the result is now fairly close to some machine instructions of the form:

```
add a, 1, 2
add b, a, 3
add c, b, 4
```

leaving the final value in register c . This suggests that this kind of transformation might be useful for compilation — more on this below.

What about functions? Suppose we define the functions

$$\begin{aligned} \text{square} &= \lambda x. x \times x \\ \text{sum_of_squares} &= \lambda x. \lambda y. (\text{square } x) + (\text{square } y) \end{aligned}$$

Applying a similar sequence of transformation as we did above to an expression like

$$\text{sum_of_squares } 10 \ 20$$

yields code that looks like:

```
let a = square 10 in
let b = square 20 in
let c = a + b in
c
```

which is okay except for the fact that *square* is actually a nontrivial expression with its own control flow kept implicit. If we want to make the control flow within *square* explicit as well, we need to do some work.

In the case of a function call, the return value of the function call is passed to the continuation that represents the rest of the program using the returned value. What if instead of just waiting for the function to return the value before passing it to the continue, we just give the continuation to the function so that it can directly call the continuation with its result?

This way, functions can be transformed into “functions that don’t return”—functions that take, besides the usual arguments, an additional argument representing a continuation. When the function finishes, it invokes the continuation on its result, instead of returning the result to its caller. Writing functions in this way is usually referred to as Continuation-Passing Style, or CPS for short.

Consider the functions above, written in such a way that they take a continuation as last argument, and invoke that continuation with a value instead of returning a value:

$$\begin{aligned} \text{square} &= \lambda x. \lambda k. k (x \times x) \\ \text{sum_of_squares} &= \lambda x. \lambda y. \lambda k. \text{square } x (\lambda a. \text{square } y (\lambda b. k (a + b))) \end{aligned}$$

Note that the last thing that code in $FACT_{cps}$ does is call a function (either k or f), and does not do anything with the result.

Continuation-Passing Style is an important concept in the compilation of functional languages and is used as an intermediate compiler representation (it has been used in compilers for Scheme, ML, etc). The main advantage is that CPS makes the control flow explicit and makes it easier to translate functional code to machine code where control is explicit (in the form of sequences of machine instructions and jumps). For instance, a CPS call can be easily translated into a jump to the invoked method, since the invoked function does not return the control.

6.1 CPS translation

We can translate lambda calculus programs into continuation-passing style. We define a translation function $\mathcal{CPS}[\cdot]$, which takes a CBV lambda calculus expression, and translates the expression to a CBV lambda calculus expression in continuation-passing style.

Let's consider a translation from lambda calculus with pairs and integers. The syntax of the source language is as follows.

$$e ::= x \mid \lambda x. e \mid e_1 e_2 \mid n \mid e_1 + e_2 \mid (e_1, e_2) \mid \#1 e \mid \#2 e$$

The translation $\mathcal{CPS}[e]$ will produce a function that whose argument is the continuation to which to pass the result. That is, for all expressions e , the translation is of the form $\mathcal{CPS}[e] = \lambda k. \dots$, where k is a continuation. We will both assume and guarantee that for any expression e , the translation $\mathcal{CPS}[e] = \lambda k. \dots$ will apply k to the result of evaluating e .

For convenience, instead of writing

$$\mathcal{CPS}[e] = \lambda k. \dots$$

we write

$$\mathcal{CPS}[e]k = \dots$$

$$\begin{aligned} \mathcal{CPS}[n]k &= k \ n \\ \mathcal{CPS}[e_1 + e_2]k &= \mathcal{CPS}[e_1] (\lambda n. \mathcal{CPS}[e_2] (\lambda m. k \ (n + m))) && n \text{ is not a free variable of } e_2 \\ \mathcal{CPS}[(e_1, e_2)]k &= \mathcal{CPS}[e_1] (\lambda v. \mathcal{CPS}[e_2] (\lambda w. k \ (v, w))) && v \text{ is not a free variable of } e_2 \\ \mathcal{CPS}[\#1 e]k &= \mathcal{CPS}[e] (\lambda v. k \ (\#1 v)) \\ \mathcal{CPS}[\#2 e]k &= \mathcal{CPS}[e] (\lambda v. k \ (\#2 v)) \\ \mathcal{CPS}[x]k &= k \ x \\ \mathcal{CPS}[\lambda x. e]k &= k \ (\lambda x. \lambda k'. \mathcal{CPS}[e]k') && k' \text{ is not a free variable of } e \\ \mathcal{CPS}[e_1 e_2]k &= \mathcal{CPS}[e_1] (\lambda f. \mathcal{CPS}[e_2] (\lambda v. f \ v \ k)) && f \text{ is not a free variable of } e_2 \end{aligned}$$

We translate a function $\lambda x. e$ to a function that takes an additional argument k' , which is the continuation after the function application. That is, k' is the continuation to which we hand the result of evaluating the function body. In function application, we see that in addition to the actual argument, we also give the continuation as the additional argument.

Let's see an example translation and execution — the initial continuation id is the identity $\lambda x. x$:

$$\begin{aligned} \mathcal{CPS}[(\lambda a. a + 6) \ 7]id &= \mathcal{CPS}[(\lambda a. a + 6)] (\lambda f. \mathcal{CPS}[7] (\lambda v. f \ v \ id)) \\ &= (\lambda f. \mathcal{CPS}[7] (\lambda v. f \ v \ id)) (\lambda a. \lambda k'. \mathcal{CPS}[a + 6]k') \\ &= (\lambda f. (\lambda v. f \ v \ id) \ 7) (\lambda a. \lambda k'. \mathcal{CPS}[a + 6]k') \\ &= (\lambda f. (\lambda v. f \ v \ id) \ 7) (\lambda a. \lambda k'. \mathcal{CPS}[a] (\lambda n. \mathcal{CPS}[6] (\lambda m. k' \ (m + n)))) \\ &= (\lambda f. (\lambda v. f \ v \ id) \ 7) (\lambda a. \lambda k'. \mathcal{CPS}[a] (\lambda n. (\lambda m. k' \ (m + n)) \ 6)) \\ &= (\lambda f. (\lambda v. f \ v \ id) \ 7) (\lambda a. \lambda k'. (\lambda n. (\lambda m. k' \ (m + n)) \ 6) \ a) \\ &\rightarrow (\lambda v. (\lambda a. \lambda k'. (\lambda n. (\lambda m. k' \ (m + n)) \ 6) \ a) \ v \ id) \ 7 \\ &\rightarrow (\lambda a. \lambda k'. (\lambda n. (\lambda m. k' \ (m + n)) \ 6) \ a) \ 7 \ id \\ &\rightarrow (\lambda n. (\lambda m. id \ (m + n)) \ 6) \ 7 \\ &\rightarrow (\lambda m. id \ (m + 7)) \ 6 \\ &\rightarrow id \ (6 + 7) \\ &\rightarrow id \ 13 \\ &\rightarrow 13 \end{aligned}$$

7 Denotational semantics of the lambda calculus

We have been treating the (pure) lambda calculus as a formal language representing functions, in some sense, but in what sense are they functions? Functions take arguments in some domain. What is the domain of the functions of the lambda calculus? Our answer has been: the functions of the lambda calculus itself. The pure lambda calculus has functions as its only values, and those functions take other functions as values. Does that even make sense?

An answer to that question would be a denotational semantics for the lambda calculus. Recall that a denotational semantics assigns to every expression in a language a mathematical object that somehow captures the behavior of that expression. Intuitively, if the lambda calculus is indeed a calculus of functions, it should be possible to give a denotational semantics to lambda terms that interprets abstractions as genuine mathematical functions.

More specifically, and somewhat loosely, is there a domain D such that for every lambda term e ,

$$\llbracket e \rrbracket \in D$$

with the property that application $e_1 e_2$ satisfies something along the lines of;

$$\llbracket e_1 e_2 \rrbracket = \llbracket e_1 \rrbracket \llbracket e_2 \rrbracket \quad ?$$

In order for such a property to hold, we need $\llbracket e_1 \rrbracket$ to really be a function $D \rightarrow D$, and therefore we need a domain D with the property that $D \rightarrow D \subseteq D$, something which cannot be by a simple counting argument. (There are way more functions in $D \rightarrow D$ than elements in D , even when D is infinite.)

Dana Scott in 1969 shows that if we restrict the space of functions to the continuous ones (for some definition of continuity), then we can in fact define such a D . Many such domains since have been discovered. Let me first show how when we have a domain D with the right properties we can give a denotational semantics to the pure lambda calculus, and then afterward give a fairly simple construction of a domain D that works.

7.1 An abstract denotational semantics

Suppose that we have a domain D that contains its own function space. In fact, we don't necessarily need it to contain its own function space, but rather, an isomorphic copy of its function space.

Technically, we assume we have a domain D with two functions φ, ψ such that

$$D \xrightleftharpoons[\psi]{\varphi} D \rightarrow D$$

where $\varphi(\psi(f)) = f$ for all $f \in D \rightarrow D$.

We will want to define the semantics inductively on the structure of lambda terms. This means that we need to be able to handle lambda terms with free variables. We will therefore define an *environment* as a map from variables to D , and define the semantics of a lambda term as a function from environments to D — intuitively, given an environment describing the value to assign to every variable in the term, the semantics will return the element of D describing the term. Formally,

$$\begin{aligned} \mathbf{Env} &= \mathbf{Var} \rightarrow D \\ \llbracket e \rrbracket &: \mathbf{Env} \rightarrow D \end{aligned}$$

As usual, for the sake of simplicity, we will write $\llbracket e \rrbracket \Gamma = \dots$ *in lieu* of $\llbracket e \rrbracket = \lambda \Gamma. \dots$

Here is the denotational semantics of the lambda calculus with respect to domain D :

$$\begin{aligned} \llbracket x \rrbracket \Gamma &= \Gamma(x) \\ \llbracket \lambda x. e \rrbracket \Gamma &= \psi(\lambda v \in D. \llbracket e \rrbracket \Gamma[x \mapsto v]) \\ \llbracket e_1 e_2 \rrbracket \Gamma &= (\varphi(\llbracket e_1 \rrbracket \Gamma)) \llbracket e_2 \rrbracket \Gamma \end{aligned}$$

Note the use of ψ and φ to convert functions in $D \rightarrow D$ into an element of D and back. Note also that we use lambda notation in the right-hand side of $\llbracket \lambda x. e \rrbracket$ to describe a mathematical function. To distinguish it more from the λ used in the syntax of the lambda calculus, I'm using a bold font for it.

Even at this juncture, there are fairly easy properties you can prove. For instance, substitution preserves denotations in the right way.

Proposition 1. *For all e_1, e_2, x, Γ , $\llbracket e_1\{e_2/x\} \rrbracket \Gamma = \llbracket e_1 \rrbracket \Gamma[x \mapsto \llbracket e_2 \rrbracket \Gamma]$.*

More importantly, the semantics is sound: β -reduction preserves denotations.

Theorem 2. *For all e_1, e_2, x , $\llbracket (\lambda x. e_1) e_2 \rrbracket = \llbracket e_1\{e_2/x\} \rrbracket$.*

Proof. Let Γ be an environment.

$$\begin{aligned} \llbracket (\lambda x. e_1) e_2 \rrbracket \Gamma &= (\varphi(\llbracket \lambda x. e_1 \rrbracket \Gamma)) \llbracket e_2 \rrbracket \Gamma \\ &= (\varphi(\psi(\lambda v \in D. \llbracket e_1 \rrbracket \Gamma[x \mapsto v]))) \llbracket e_2 \rrbracket \Gamma \\ &= (\lambda v \in D. \llbracket e_1 \rrbracket \Gamma[x \mapsto v]) \llbracket e_2 \rrbracket \Gamma \\ &= \llbracket e_1 \rrbracket \Gamma[x \mapsto \llbracket e_2 \rrbracket \Gamma] \\ &= \llbracket e_1\{e_2/x\} \rrbracket \Gamma \end{aligned}$$

□

Note where we use the fact that $\varphi(\psi(f)) = f$ in the proof.

An easy induction yields that whenever $e \rightarrow^* e'$, then $\llbracket e \rrbracket = \llbracket e' \rrbracket$.

7.2 Existence of a model

The above shows that if we have a suitable D , then we can give a semantics to the lambda calculus. Does such a D exist? Dana Scott in 1969 exhibited a domain D_∞ that did the job, but the construction is abstract enough that it is difficult to visualize. (It is obtained by a limiting process, by constructing a sequence of domains D_0, D_1, D_2, \dots) In 1972, Gordon Plotkin discovered a more concrete model \mathcal{P}_ω that we give here.

The domain \mathcal{P}_ω is a CPO (in fact, a complete lattice, where all sets have least upperbounds, not just ω -chains), and getting an isomorphic copy of its function space inside \mathcal{P}_ω is achieved by considering only the *continuous functions* from \mathcal{P}_ω to \mathcal{P}_ω , written $[\mathcal{P}_\omega \rightarrow \mathcal{P}_\omega]$. Recall that a function over a partial order is continuous if it preserves least upperbounds, that is, if $f(\sqcup X) = \sqcup \{f(x) \mid x \in X\}$

The domain \mathcal{P}_ω is simply the collection $\{X \mid X \subseteq \mathbb{N}\}$ of all subsets of \mathbb{N} , ordered by \subseteq .

Proposition 2.

(a) *For all $X \in \mathcal{P}_\omega$, $X = \bigcup \{S \mid S \text{ finite and } S \subseteq X\}$*

(b) *A function $f : \mathcal{P}_\omega \rightarrow \mathcal{P}_\omega$ is continuous if and only if for all $X \in \mathcal{P}_\omega$,*

$$f(X) = \bigcup \{f(S) \mid S \text{ finite and } S \subseteq X\}$$

Part (b) says that a continuous function is characterized by its behavior on the finite subsets of \mathbb{N} . In other words, to fully describe a continuous function $\mathcal{P}_\omega \rightarrow \mathcal{P}_\omega$, it suffices to state what it maps the finite subsets of \mathbb{N} to.

There are only countably many finite subsets of \mathbb{N} , and we can enumerate them. Define

$$e_n = \{k_1, \dots, k_m\} \quad \text{for } k_1 < \dots < k_m \text{ and } n = \sum_{i=1}^m 2^{k_i}$$

In other words, the n th set is the set of positions in the binary representation of n that have a 1. Clearly, every finite subset of \mathbb{N} is e_n for some n , and every e_n is distinct and finite for different values of n . Thus, e_0, e_1, e_2, \dots is an enumeration of all the finite subsets of \mathbb{N} .

This means that $f : \mathcal{P}_\omega \longrightarrow \mathcal{P}_\omega$ is continuous if and only if $f(X) = \bigcup \{f(e_n) \mid e_n \subseteq X\}$. Since $f(e_n)$ is a subset of \mathbb{N} , to fully describe f , it suffices to describe, for every n , every natural number m in $f(e_n)$. There are countably many such numbers for every n , and therefore there are countably many pairs of m and n needed to describe f .

We can encode a pair of natural numbers into a single natural in such a way that we can recover the initial pair uniquely. Here is such an encoding:

$$(m, n) = \frac{1}{2}(m+n)(m+n+1) + n$$

This yields $(0, 0) = 0$, $(1, 0) = 1$, $(0, 1) = 2$, $(2, 0) = 3$, $(1, 1) = 4$, $(0, 2) = 5$, \dots

I claim that we can represent a continuous function $f : \mathcal{P}_\omega \longrightarrow \mathcal{P}_\omega$ as the set $\{(m, n) \mid m \in f(e_n)\}$, and this is a subset of \mathbb{N} via the encoding of pairs given above. This gives us a function

$$\text{graph}(f) = \{(m, n) \mid m \in f(e_n)\}$$

from $[\mathcal{P}_\omega \longrightarrow \mathcal{P}_\omega] \longrightarrow \mathcal{P}_\omega$. Given the “graph” of a function, we can recover the original function via:

$$\text{fun}(X) = \lambda V \in \mathcal{P}_\omega. \{m \mid \exists e_n \subseteq V. (m, n) \in X\}$$

which is a function $\mathcal{P}_\omega \longrightarrow [\mathcal{P}_\omega \longrightarrow \mathcal{P}_\omega]$. Therefore, we have

$$\mathcal{P}_\omega \xrightleftharpoons[\text{graph}]{\text{fun}} [\mathcal{P}_\omega \longrightarrow \mathcal{P}_\omega]$$

And it is reasonably easy to check that

$$\text{fun}(\text{graph}(f)) = f$$

which is the property we need for \mathcal{P}_ω to be a suitable model for the lambda calculus.¹

¹Technically, we also need to show that the functions used in the denotational semantics using \mathcal{P}_ω are continuous, since the function space used is that of continuous functions. That argument it turns out can be made generically for all CPOs D .