

Notes on Typed Lambda Calculi

Spring 2017

A *type* is a collection of computational entities that share some common property. For example, the type **int** represents all expressions that evaluate to an integer, and the type **int** \rightarrow **int** represents all functions from integers to integers. The Pascal subrange type $[1..100]$ represents all integers between 1 and 100.

Types can be thought of as describing computations succinctly and approximately: types are a *static* approximation to the run-time behaviors of terms and programs. Type systems are a lightweight formal method for reasoning about behavior of a program. Uses of type systems include: naming and organizing useful concepts; providing information (to the compiler or programmer) about data manipulated by a program; and ensuring that the run-time behavior of programs meet certain criteria.

In this lecture, we'll consider a type system for the lambda calculus that ensures that values are used correctly; for example, that a program never tries to add an integer to a function. The resulting language (lambda calculus plus the type system) is called the *simply-typed lambda calculus*.

1 Simply-typed lambda calculus

The syntax of the simply-typed lambda calculus is similar to that of untyped lambda calculus, with the exception of abstractions. Since abstractions define functions that take an argument, in the simply-typed lambda calculus, we explicitly state what the type of the argument is. That is, in an abstraction $\lambda x:\tau. e$, the τ is the expected type of the argument.

The syntax of the simply-typed lambda calculus is as follows. We will include integer literals n , addition $e_1 + e_2$, and the *unit value* $()$. The unit value is the only value of type **unit**.

expressions	$e ::= x \mid \lambda x:\tau. e \mid e_1 \ e_2 \mid n \mid e_1 + e_2 \mid ()$
values	$v ::= \lambda x:\tau. e \mid n \mid ()$
types	$\tau ::= \mathbf{int} \mid \mathbf{unit} \mid \tau_1 \rightarrow \tau_2$

The operational semantics of the simply-typed lambda calculus are the same as the untyped lambda calculus. For completeness, we present the CBV small step operational semantics here.

$$\begin{array}{c}
 E ::= [\cdot] \mid E \ e \mid v \ E \mid E + e \mid v + E \\
 \text{CONTEXT} \frac{e \rightarrow e'}{E[e] \rightarrow E[e']} \\
 \beta\text{-REDUCTION} \frac{}{(\lambda x. e) \ v \rightarrow e\{v/x\}} \qquad \text{ADD} \frac{}{n_1 + n_2 \rightarrow n} n = n_1 + n_2
 \end{array}$$

1.1 The typing relation

The presence of types does not alter the evaluation of an expression at all. So what use are types?

We will use types to restrict what expressions we will evaluate. Specifically, the type system for the simply-typed lambda calculus will ensure that any *well-typed* program will not get *stuck*. A term e is stuck if e is not a value and there is no term e' such that $e \rightarrow e'$. For example, the expression $42 + \lambda x:\mathbf{int}. x$ is stuck: it attempts to add an integer and a function; it is not a value, and there is no operational rule that allows us to reduce this expression. Another stuck expression is $() \ 47$, which attempts to apply the unit value to an integer.

We introduce a relation (or *judgment*) over *typing contexts* (or *type environments*) Γ , expressions e , and types τ . The judgment

$$\Gamma \vdash e:\tau$$

is read as “ e has type τ in context Γ ”.

A typing context is a sequence of variables and their types. In the typing judgment $\Gamma \vdash e : \tau$, we will ensure that if x is a free variable of e , then Γ associates x with a type. We can view a typing context as a partial function from variables to types. We will write $\Gamma, x : \tau$ or $\Gamma[x \mapsto \tau]$ to indicate the typing context that extends Γ by associating variable x with type τ . The empty context is sometimes written \emptyset , or often just not written at all. For example, we write $\vdash e : \tau$ to mean that the closed term e has type τ under the empty context.

Given a typing environment Γ and expression e , if there is some τ such that $\Gamma \vdash e : \tau$, we say that e is *well-typed under context* Γ ; if Γ is the empty context, we say e is *well-typed*.

We define the judgment $\Gamma \vdash e : \tau$ inductively.

$$\begin{array}{c} \text{T-INT} \frac{}{\Gamma \vdash n : \mathbf{int}} \quad \text{T-ADD} \frac{\Gamma \vdash e_1 : \mathbf{int} \quad \Gamma \vdash e_2 : \mathbf{int}}{\Gamma \vdash e_1 + e_2 : \mathbf{int}} \quad \text{T-UNIT} \frac{}{\Gamma \vdash () : \mathbf{unit}} \\[10pt] \text{T-VAR} \frac{}{\Gamma \vdash x : \tau} \quad \Gamma(x) = \tau \quad \text{T-ABS} \frac{\Gamma, x : \tau \vdash e : \tau'}{\Gamma \vdash \lambda x : \tau. e : \tau \rightarrow \tau'} \quad \text{T-APP} \frac{\Gamma \vdash e_1 : \tau \rightarrow \tau' \quad \Gamma \vdash e_2 : \tau}{\Gamma \vdash e_1 e_2 : \tau'} \end{array}$$

An integer n always has type **int**. Expression $e_1 + e_2$ has type **int** if both e_1 and e_2 have type **int**. The unit value $()$ always has type **unit**.

Variable x has whatever type the context associates with x . Note that Γ must contain an associating for x in order to the judgment $\Gamma \vdash x : \tau$ to hold, that is, $x \in \text{dom}(\Gamma)$. The abstraction $\lambda x : \tau. e$ has the function type $\tau \rightarrow \tau'$ if the function body e has type τ' under the assumption that x has type τ . Finally, an application $e_1 e_2$ has type τ' provided that e_1 is a function of type $\tau \rightarrow \tau'$, and e_2 is an argument of the expected type, i.e., of type τ .

To type check an expression e , we attempt to construct a derivation of the judgment $\vdash e : \tau$, for some type τ . For example, consider the program $(\lambda x : \mathbf{int}. x + 40) 2$. The following is a proof that $(\lambda x : \mathbf{int}. x + 40) 2$ is well-typed.

$$\begin{array}{c} \text{T-VAR} \frac{}{x : \mathbf{int} \vdash x : \mathbf{int}} \quad \text{T-INT} \frac{}{x : \mathbf{int} \vdash 40 : \mathbf{int}} \\ \text{T-ADD} \frac{}{x : \mathbf{int} \vdash x + 40 : \mathbf{int}} \\ \text{T-ABS} \frac{}{\vdash \lambda x : \mathbf{int}. x + 40 : \mathbf{int} \rightarrow \mathbf{int}} \\ \text{T-APP} \frac{}{\vdash (\lambda x : \mathbf{int}. x + 40) 2 : \mathbf{int}} \quad \text{T-INT} \frac{}{\vdash 2 : \mathbf{int}} \end{array}$$

1.2 Type soundness

We mentioned above that the type system ensures that any well-typed program does not get stuck. We can state this property formally.

Theorem (Type soundness). *If $\vdash e : \tau$ and $e \longrightarrow^* e'$ then either e' is a value, or there exists e'' such that $e' \longrightarrow e''$.*

We will prove this theorem using two lemmas: *preservation* and *progress*. Intuitively, preservation says that if an expression e is well-typed, and e can take a step to e' , then e' is well-typed. That is, evaluation preserves well-typedness. Progress says that if an expression e is well-typed, then either e is a value, or there is an e' such that e can take a step to e' . That is, well-typedness means that the expression cannot get stuck. Together, these two lemmas suffice to prove type soundness.

1.2.1 Preservation

Lemma (Preservation). *If $\vdash e : \tau$ and $e \longrightarrow e'$ then $\vdash e' : \tau$.*

Proof. We proceed by induction on $e \longrightarrow e'$. That is, we will prove for all e and e' such that $e \longrightarrow e'$, that $P(e \longrightarrow e')$ holds, where

$$P(e \longrightarrow e') = \forall \tau. \text{ if } \vdash e : \tau \text{ then } \vdash e' : \tau.$$

Consider each of the inference rules for the small step relation.

- ADD

Assume $\vdash e : \tau$.

Here $e \equiv n_1 + n_2$, and $e' = n$ where $n = n_1 + n_2$, and $\tau = \mathbf{int}$. By the typing rule T-INT, we have $\vdash e' : \mathbf{int}$ as required.

- β -REDUCTION

Assume $\vdash e : \tau$.

Here, $e \equiv (\lambda x : \tau'. e_1) v$ and $e' \equiv e_1\{v/x\}$. Since e is well-typed, we have derivations showing $\vdash \lambda x : \tau'. e_1 : \tau' \rightarrow \tau$ and $\vdash v : \tau'$. There is only one typing rule for abstractions, T-ABS, from which we know $x : \tau \vdash e_1 : \tau$. By the substitution lemma (see below), we have $\vdash e_1\{v/x\} : \tau$ as required.

- CONTEXT

Assume $\vdash e : \tau$.

Here, we have some context E such that $e = E[e_1]$ and $e' = E[e_2]$ for some e_1 and e_2 such that $e_1 \longrightarrow e_2$. The inductive hypothesis is that $P(e_1 \longrightarrow e_2)$.

Since e is well-typed, we can show by induction on the structure of E that $\vdash e_1 : \tau_1$ for some τ_1 . By the inductive hypothesis, we thus have $\vdash e_2 : \tau_1$. By the context lemma (see below) we have $\vdash E[e'] : \tau$ as required.

□

Additional lemmas we used in the proof above.

Lemma (Substitution). *If $x : \tau' \vdash e : \tau$ and $\vdash v : \tau'$ then $\vdash e\{v/x\} : \tau$.*

Lemma (Context). *If $\vdash E[e_0] : \tau$ and $\vdash e_0 : \tau'$ and $\vdash e_1 : \tau'$ then $\vdash E[e_1] : \tau$.*

1.2.2 Progress

Lemma (Progress). *If $\vdash e : \tau$ then either e is a value or there exists an e' such that $e \longrightarrow e'$.*

Proof. We proceed by induction on the derivation of $\vdash e : \tau$. That is, we will show for all e and τ such that $\vdash e : \tau$, we have $P(\vdash e : \tau)$, where

$$P(\vdash e : \tau) = \text{either } e \text{ is a value or } \exists e' \text{ such that } e \longrightarrow e'.$$

- T-VAR

This case is impossible, since a variable is not well-typed in the empty environment.

- T-UNIT, T-INT, T-ABS

Trivial, since e must be a value.

- T-ADD

Here $e \equiv e_1 + e_2$ and $\vdash e_i : \mathbf{int}$ for $i \in \{1, 2\}$. By the inductive hypothesis, for $i \in \{1, 2\}$, either e_i is a value or there is an e'_i such that $e_i \longrightarrow e'_i$.

If e_1 is not a value, then by CONTEXT, $e_1 + e_2 \longrightarrow e'_1 + e_2$. If e_1 is a value and e_2 is not a value, then by CONTEXT, $e_1 + e_2 \longrightarrow e_1 + e'_2$. If e_1 and e_2 are values, then, it must be the case that they are both integer literals, and so, by ADD, we have $e_1 + e_2 \longrightarrow n$ where n equals e_1 plus e_2 .

- T-APP

Here $e \equiv e_1 e_2$ and $\vdash e_1 : \tau' \rightarrow \tau$ and $\vdash e_2 : \tau'$. By the inductive hypothesis, for $i \in \{1, 2\}$, either e_i is a value or there is an e'_i such that $e_i \longrightarrow e'_i$.

If e_1 is not a value, then by CONTEXT, $e_1 e_2 \longrightarrow e'_1 e_2$. If e_1 is a value and e_2 is not a value, then by CONTEXT, $e_1 e_2 \longrightarrow e_1 e'_2$. If e_1 and e_2 are values, then, it must be the case that e_1 is an abstraction $\lambda x : \tau'. e'$, and so, by β -REDUCTION, we have $e_1 e_2 \longrightarrow e' \{e_2/x\}$.

□

1.3 Expressive power of the simply-typed lambda calculus

Clearly, not all expressions in the untyped lambda calculus are well-typed. Indeed, type soundness implies that any lambda calculus program that gets stuck is not well-typed. But are there programs that do not get stuck that are not well-typed?

Unfortunately, the answer is yes.

First, since the simply-typed lambda calculus requires us to specify a type for function arguments, any given function can only take arguments of one type. Consider, for example, the identity function $\lambda x. x$. This function may be applied to any argument, and it will not get stuck. However, we must provide a type for the argument. If we specify $\lambda x : \mathbf{int}. x$, then this function can only accept integers, and the program $(\lambda x : \mathbf{int}. x) ()$ is not well-typed, even though it does not get stuck. Indeed, in the simply-typed lambda calculus, there is a different identity function for each type.

Second, we can no longer write recursive functions. Consider the nonterminating expression $\Omega = (\lambda x. x x) (\lambda x. x x)$. What type does it have? Let's suppose that the type of $\lambda x. x x$ is $\tau \rightarrow \tau'$. But $\lambda x. x x$ is applied to itself! So that means that the type of $\lambda x. x x$ is the argument type τ . So we have that τ must be equal to $\tau \rightarrow \tau'$. There is no such type for which this equality holds. (At least, not in this type system...)

This means that every well-typed program in the simply-typed lambda calculus terminates. More formally:

Theorem 1 (Normalization). *If $\vdash e : \tau$ then there exists a value v such that $e \longrightarrow^* v$.*

This is known as *normalization* since it means that given any well-typed expression, we can reduce it to a *normal form*, which, in our case, is a value.

2 Products and sums

We have previously seen *products*, which are pairs of expressions. Products were constructed using the expression (e_1, e_2) , and destructured using projection $\#1 e$ and $\#2 e$.

$$\begin{aligned} e &::= \dots \mid (e_1, e_2) \mid \#1 e \mid \#2 e \\ v &::= \dots \mid (v_1, v_2) \end{aligned}$$

Again, there are structural rules to determine the order of evaluation. In a CBV lambda calculus, the evaluation contexts are extended as follows.

$$E ::= \dots \mid (E, e) \mid (v, E) \mid \#1 E \mid \#2 E$$

In addition to the structural rules, there are two operational semantics rules that show how the destructors and constructor interact.

$$\frac{}{\#1 (v_1, v_2) \longrightarrow v_1} \qquad \frac{}{\#2 (v_1, v_2) \longrightarrow v_2}$$

The type of a product expression (or a *product type*) is a pair of types, written $\tau_1 \times \tau_2$. The typing rules for the product constructors and destructors are the following.

$$\frac{\Gamma \vdash e_1 : \tau_1 \quad \Gamma \vdash e_2 : \tau_2}{\Gamma \vdash (e_1, e_2) : \tau_1 \times \tau_2}$$

$$\frac{\Gamma \vdash e : \tau_1 \times \tau_2}{\Gamma \vdash \#1 \ e : \tau_1}$$

$$\frac{\Gamma \vdash e : \tau_1 \times \tau_2}{\Gamma \vdash \#2 \ e : \tau_2}$$

We introduce *sums*, which are dual to products. Intuitively, a product holds two values, one of type τ_1 , and one of type τ_2 . By contrast, a sum holds a single value that is either of type τ_1 or of type τ_2 . The type of a sum is written $\tau_1 + \tau_2$. There are two constructors for a sum, corresponding to whether we are constructing a sum with a value of τ_1 or a value of τ_2 .

$$\begin{aligned} e &::= \dots \mid \text{inl}_{\tau_1 + \tau_2} \ e \mid \text{inr}_{\tau_1 + \tau_2} \ e \mid \text{case } e_1 \text{ of } e_2 \mid e_3 \\ v &::= \dots \mid \text{inl}_{\tau_1 + \tau_2} \ v \mid \text{inr}_{\tau_1 + \tau_2} \ v \end{aligned}$$

Again, there are structural rules to determine the order of evaluation. In a CBV lambda calculus, the evaluation contexts are extended as follows.

$$E ::= \dots \mid \text{inl}_{\tau_1 + \tau_2} \ E \mid \text{inr}_{\tau_1 + \tau_2} \ E \mid \text{case } E \text{ of } e_2 \mid e_3$$

In addition to the structural rules, there are two operational semantics rules that show how the destructors and constructors interact.

$$\frac{}{\text{case inl}_{\tau_1 + \tau_2} \ v \text{ of } e_2 \mid e_3 \longrightarrow e_2 \ v} \quad \frac{}{\text{case inr}_{\tau_1 + \tau_2} \ v \text{ of } e_2 \mid e_3 \longrightarrow e_3 \ v}$$

The type of a sum expression (or a *sum type*) is written $\tau_1 + \tau_2$. The typing rules for the sum constructors and destructor are the following.

$$\frac{\Gamma \vdash e : \tau_1}{\Gamma \vdash \text{inl}_{\tau_1 + \tau_2} \ e : \tau_1 + \tau_2} \quad \frac{\Gamma \vdash e : \tau_2}{\Gamma \vdash \text{inr}_{\tau_1 + \tau_2} \ e : \tau_1 + \tau_2} \quad \frac{\Gamma \vdash e : \tau_1 + \tau_2 \quad \Gamma \vdash e_1 : \tau_1 \rightarrow \tau \quad \Gamma \vdash e_2 : \tau_2 \rightarrow \tau}{\Gamma \vdash \text{case } e \text{ of } e_1 \mid e_2 : \tau}$$

Let's see an example of a program that uses sum types.

```
let f : (int + (int → int)) → int =
  λa : int + (int → int). case a of λy. y + 1 | λg. g 35 in
let h : int → int = λx : int. x + 7 in
f (inrint+(int→int) h)
```

Here, the function f takes argument a , which is a sum. That is, the actual argument for a will either be a value of type **int** or a value of type **int** \rightarrow **int**. We destroy the sum value with a case statement, which must be prepared to take either of the two kinds of values that the sum may contain. We end up applying f to a value of type **int** \rightarrow **int** (i.e., a value injected into the right type of the sum). The entire program ends up evaluating to 42.

3 Recursion

We saw in last lecture that we could not type recursive functions or fixed-point combinators in the simply-typed lambda calculus. So instead of trying (and failing) to define a fixed-point combinator in the simply-typed lambda calculus, we add a new primitive $\mu x : \tau. e$ to the language. The evaluation rules for the new primitive will mimic the behavior of fixed-point combinators.

We extend the syntax with the new primitive operator. Intuitively, $\mu x : \tau. e$ is the fixed-point of the function $\lambda x : \tau. e$. Note that $\mu x : \tau. e$ is *not* a value, regardless of whether e is a value or not.

$$e ::= \dots \mid \mu x : \tau. e$$

We extend the operational semantics for the new operator. There is a new axiom, but no new evaluation contexts.

$$\overline{\mu x:\tau. e \longrightarrow e\{(\mu x:\tau. e)/x\}}$$

Note that we can define the $\text{letrec } x:\tau = e_1 \text{ in } e_2$ construct in terms of this new expression.

$$\text{letrec } x:\tau = e_1 \text{ in } e_2 \triangleq \text{let } x:\tau = \mu x:\tau. e_1 \text{ in } e_2$$

We add a new typing rule for the new language construct.

$$\frac{\Gamma[x \mapsto \tau] \vdash e:\tau}{\Gamma \vdash \mu x:\tau. e:\tau}$$

Returning to our trusty factorial example, the following program implements the factorial function using the $\mu x:\tau. e$ expression.

$$FACT \triangleq \mu f:\mathbf{int} \rightarrow \mathbf{int}. \lambda n:\mathbf{int}. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f \ (n - 1))$$

Or using our convenient letrec notation, we could define a variable $fact$ as follows.

$$\begin{aligned} \text{letrec } fact:\mathbf{int} \rightarrow \mathbf{int} = \lambda n:\mathbf{int}. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (fact \ (n - 1)) \\ \text{in } \dots \end{aligned}$$

We can write non-terminating computations for any type: the expression $\mu x:\tau. x$ has type τ , and does not terminate.

Although the $\mu x:\tau. e$ expression is normally used to define recursive functions, it can be used to find fixed points of any type. For example, consider the following expression.

$$\begin{aligned} \mu x:(\mathbf{int} \rightarrow \mathbf{bool}) \times (\mathbf{int} \rightarrow \mathbf{bool}). (\lambda n:\mathbf{int}. \text{if } n = 0 \text{ then true else } ((\#2 \ x) \ (n - 1))), \\ \lambda n:\mathbf{int}. \text{if } n = 0 \text{ then false else } ((\#1 \ x) \ (n - 1))) \end{aligned}$$

This expression has type $(\mathbf{int} \rightarrow \mathbf{bool}) \times (\mathbf{int} \rightarrow \mathbf{bool})$ —it is a pair of mutually recursive functions; the first function returns **true** only if its argument is even; the second returns **true** only if its argument is odd.

4 Denotational semantics

The denotational semantics of the simply-typed lambda calculus is much more straightforward than that of the untyped lambda calculus, because types provide helpful structure.

We first consider the CBV simply-typed lambda calculus with products and sums:

$$\begin{aligned} e &::= x \mid \lambda x:\tau. e \mid e_1 \ e_2 \mid n \mid e_1 + e_2 \mid () \mid (e_1, e_2) \mid \#1 \ e \mid \#2 \ e \mid \text{inl}_{\tau_1+\tau_2} \ e \mid \text{inr}_{\tau_1+\tau_2} \ e \mid \text{case } e_1 \text{ of } e_2 \mid e_3 \\ v &::= \lambda x:\tau. e \mid n \mid () \mid (v_1, v_2) \mid \text{inl}_{\tau_1+\tau_2} \ v \mid \text{inr}_{\tau_1+\tau_2} \ v \\ \tau &::= \mathbf{int} \mid \mathbf{unit} \mid \tau_1 \rightarrow \tau_2 \mid \tau_1 \times \tau_2 \mid \tau_1 + \tau_2 \end{aligned}$$

From the operational semantics, we know that well-typed terms reduce to values (by strong normalization) of the same type as the original expression (by type soundness). Thus, a reasonable denotational semantics is to associate with every term the resulting value to which it reduces. Thus, the semantics associates with terms of type τ results of type τ . The map $\llbracket \tau \rrbracket$ describes the domain of type τ :

$$\begin{aligned} \llbracket \mathbf{int} \rrbracket &= \mathbb{Z} \\ \llbracket \mathbf{unit} \rrbracket &= \{\bullet\} \\ \llbracket \tau_1 \rightarrow \tau_2 \rrbracket &= \llbracket \tau_1 \rrbracket \rightarrow \llbracket \tau_2 \rrbracket \\ \llbracket \tau_1 \times \tau_2 \rrbracket &= \llbracket \tau_1 \rrbracket \times \llbracket \tau_2 \rrbracket \\ \llbracket \tau_1 + \tau_2 \rrbracket &= \llbracket \tau_1 \rrbracket \uplus \llbracket \tau_2 \rrbracket \end{aligned}$$

where for the sake of concreteness we take $X \uplus Y = \{(0, x) \mid x \in X\} \cup \{(1, y) \mid y \in Y\}$.

Since we only care about well-typed terms, we only give a semantics to such. We will do so by assigning a meaning to *typing judgments*: we define a semantics $\llbracket \Gamma \vdash e : \tau \rrbracket$. Since the semantics is defined in terms of a typing derivation, we have the subtrees of the derivation available to define the semantics of the subterms.

To account for open terms (since subterms of even a closed terms may be open), the meaning of a term e of type τ will be a function from environments (giving a value for every variable) to values in $\llbracket \tau \rrbracket$.

An environment ρ is a map

$$\mathbf{Var} \rightarrow \bigcup \{ \llbracket \tau \rrbracket \mid \tau \text{ a type} \}$$

assigning to every variable a value (of any type). Let \mathbf{Env}_Γ be the set of environments that respect the typing context Γ : the set of all ρ such that for all $x \in \text{Dom}(\Gamma)$, $\rho(x) \in \llbracket \Gamma(\tau) \rrbracket$.

$$\llbracket \Gamma \vdash e : \tau \rrbracket : \mathbf{Env}_\Gamma \rightarrow \llbracket \tau \rrbracket$$

$$\begin{aligned} \llbracket \Gamma \vdash x : \tau \rrbracket \rho &= \rho(x) \\ \llbracket \Gamma \vdash \lambda x : \tau. e : \tau \rightarrow \tau' \rrbracket \rho &= \lambda v \in \llbracket \tau \rrbracket. \llbracket \Gamma, x : \tau \vdash e : \tau' \rrbracket \rho[x \mapsto v] \\ \llbracket \Gamma \vdash e_1 \ e_2 : \tau \rrbracket \rho &= (\llbracket \Gamma \vdash e_1 : \tau' \rightarrow \tau \rrbracket \rho) (\llbracket \Gamma \vdash e_2 : \tau' \rrbracket \rho) \\ \llbracket \Gamma \vdash n : \mathbf{int} \rrbracket \rho &= n \\ \llbracket \Gamma \vdash e_1 + e_2 : \mathbf{int} \rrbracket \rho &= \llbracket \Gamma \vdash e_1 : \mathbf{int} \rrbracket \rho + \llbracket \Gamma \vdash e_2 : \mathbf{int} \rrbracket \rho \\ \llbracket \Gamma \vdash () : \mathbf{unit} \rrbracket \rho &= \bullet \\ \llbracket \Gamma \vdash (e_1, e_2) : \tau_1 \times \tau_2 \rrbracket \rho &= (\llbracket \Gamma \vdash e_1 : \tau_1 \rrbracket \rho, \llbracket \Gamma \vdash e_2 : \tau_2 \rrbracket \rho) \\ \llbracket \Gamma \vdash \#1 \ e : \tau_1 \rrbracket \rho &= \pi_1(\llbracket \Gamma \vdash e : \tau_1 \times \tau_2 \rrbracket \rho) \\ \llbracket \Gamma \vdash \#2 \ e : \tau_2 \rrbracket \rho &= \pi_2(\llbracket \Gamma \vdash e : \tau_1 \times \tau_2 \rrbracket \rho) \\ \llbracket \Gamma \vdash \mathbf{inl}_{\tau_1 + \tau_2} \ e : \tau_1 + \tau_2 \rrbracket \rho &= (0, \llbracket \Gamma \vdash e : \tau_1 \rrbracket \rho) \\ \llbracket \Gamma \vdash \mathbf{inr}_{\tau_1 + \tau_2} \ e : \tau_1 + \tau_2 \rrbracket \rho &= (1, \llbracket \Gamma \vdash e : \tau_2 \rrbracket \rho) \\ \llbracket \Gamma \vdash \mathbf{case} \ e_1 \ \mathbf{of} \ e_2 \mid e_3 : \tau \rrbracket \rho &= (\llbracket \Gamma \vdash e_2 : \tau_1 \rightarrow \tau \rrbracket \rho, \llbracket \Gamma \vdash e_3 : \tau_2 \rightarrow \tau \rrbracket \rho) (\llbracket \Gamma \vdash e_1 : \tau_1 + \tau_2 \rrbracket \rho) \end{aligned}$$

where $[f, g]$ is defined by:

$$[f, g] (i, v) = \begin{cases} f(v) & \text{if } i = 0 \\ g(v) & \text{if } i = 1 \end{cases}$$

Thus, roughly, functions in the simply-typed lambda calculus correspond to functions over type domains, products in the simply-typed lambda calculus correspond to Cartesian products over type domains, and sums in the simply-typed lambda calculus correspond to disjoin unions over type domains. All in all, a fairly natural correspondence.

To what extent does the above denotational semantics capture the operational semantics? The following result is reasonably easy to prove, by structural induction on the reduction relation \longrightarrow :

Theorem 2. *If $\vdash e : \tau$ and $e \longrightarrow e'$, then $\llbracket \vdash e : \tau \rrbracket = \llbracket \vdash e' : \tau \rrbracket$.*

An easy induction then yields

Corollary 1. *If $\vdash e : \tau$ and $e \longrightarrow^* v$, then $\llbracket \vdash e : \tau \rrbracket = \llbracket \vdash v : \tau \rrbracket$.*

In other words, the denotational semantics of a term is preserved by the operational semantics. What about the other direction? If an expression and a value have the same denotation, does the expression reduce to the value? Unfortunately, that result does *not* hold in general. Consider $e = \lambda x : \mathbf{int}. x + 0$ and $v = \lambda x : \mathbf{int}. x$. It is easy to see that $\llbracket \vdash e : \mathbf{int} \rightarrow \mathbf{int} \rrbracket = \llbracket \vdash v : \mathbf{int} \rightarrow \mathbf{int} \rrbracket = \lambda v \in \mathbb{Z}.v$, but we cannot have $e \longrightarrow^* v$ since e is already a value and does not reduce any further.

What we *can* show is that as long as type τ does not contain a function type, then the converse to Corollary 1 holds.

Theorem 3. *If $\vdash e : \tau$ for τ not containing a function type and $\llbracket \vdash e : \tau \rrbracket = \llbracket \vdash v : \tau \rrbracket$, then $e \longrightarrow^* v$.*

The proof of this theorem is a bit more involved, though, and requires the use of a new technique called *logical relations*. (See Chapter 11 of Winskel's *The Formal Semantics of Programming Languages*, for example.)

What about recursion? If we add $\mu x:\tau. e$ to the language, how can we adapt the denotational semantics? Just like we did for the denotational semantics of IMP, one way is to move to partial functions.

To simplify the presentation somewhat, we first restrict recursion so that in $\mu x:\tau. e$ expression e is actually an abstraction $\lambda y:\tau_1. e$. (In a call-by-value language, the only useful recursive artifacts are recursive functions.) Thus, we restrict recursive expressions to $\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e$.

The domain of type $\tau_1 \rightarrow \tau_2$ is now the set $\llbracket \tau_1 \rrbracket \rightarrow \llbracket \tau_2 \rrbracket$ of partial functions from $\llbracket \tau_1 \rrbracket$ to $\llbracket \tau_2 \rrbracket$, and the denotation of a judgment $\Gamma \vdash e:\tau$ is now a partial function from environments to the domain $\llbracket \tau \rrbracket$:

$$\llbracket \Gamma \vdash e:\tau \rrbracket : \mathbf{Env}_\Gamma \rightarrow \llbracket \tau \rrbracket$$

The denotational semantics are now an adaptation of the above denotational to partial functions. As we have done in the past, we freely use the graph representation of partial functions to simplify the definition of partial functions.

$$\begin{aligned} \llbracket \Gamma \vdash x:\tau \rrbracket &= \{(\rho, \rho(x)) \mid \rho \in \mathbf{Env}_\Gamma\} \\ \llbracket \Gamma \vdash \lambda x:\tau. e:\tau \rightarrow \tau' \rrbracket &= \{(\rho, f) \mid \rho \in \mathbf{Env}_\Gamma, f = \{(v, v') \mid (\rho[x \mapsto v], v') \in \llbracket \Gamma, x:\tau \vdash e:\tau' \rrbracket\}\} \\ \llbracket \Gamma \vdash e_1 e_2:\tau \rrbracket &= \{(\rho, v') \mid \rho \in \mathbf{Env}_\Gamma, (\rho, f) \in \llbracket \Gamma \vdash e_1:\tau' \rightarrow \tau \rrbracket, (\rho, v) \in \llbracket \Gamma \vdash e_2:\tau' \rrbracket, (v, v') \in f\} \\ \llbracket \Gamma \vdash n:\mathbf{int} \rrbracket &= \{(\rho, n) \mid \rho \in \mathbf{Env}_\Gamma\} \\ \llbracket \Gamma \vdash e_1 + e_2:\mathbf{int} \rrbracket &= \{(\rho, n) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, n_1) \in \llbracket \Gamma \vdash e_1:\mathbf{int} \rrbracket, (\rho, n_2) \in \llbracket \Gamma \vdash e_2:\mathbf{int} \rrbracket, n = n_1 + n_2\} \\ \llbracket \Gamma \vdash ():\mathbf{unit} \rrbracket &= \{(\rho, \bullet) \mid \rho \in \mathbf{Env}_\Gamma\} \\ \llbracket \Gamma \vdash (e_1, e_2):\tau_1 \times \tau_2 \rrbracket &= \{(\rho, (v_1, v_2)) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, v_1) \in \llbracket \Gamma \vdash e_1:\tau_1 \rrbracket, (\rho, v_2) \in \llbracket \Gamma \vdash e_2:\tau_2 \rrbracket\} \\ \llbracket \Gamma \vdash \#1 e:\tau_1 \rrbracket &= \{(\rho, v_1) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, (v_1, v_2)) \in \llbracket \Gamma \vdash e:\tau_1 \times \tau_2 \rrbracket\} \\ \llbracket \Gamma \vdash \#2 e:\tau_2 \rrbracket &= \{(\rho, v_2) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, (v_1, v_2)) \in \llbracket \Gamma \vdash e:\tau_1 \times \tau_2 \rrbracket\} \\ \llbracket \Gamma \vdash \mathbf{inl}_{\tau_1+\tau_2} e:\tau_1 + \tau_2 \rrbracket &= \{(\rho, (0, v)) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, v) \in \llbracket \Gamma \vdash e:\tau_1 \rrbracket\} \\ \llbracket \Gamma \vdash \mathbf{inr}_{\tau_1+\tau_2} e:\tau_1 + \tau_2 \rrbracket &= \{(\rho, (1, v)) \mid \rho \in \mathbf{Env}_\Gamma, (\rho, v) \in \llbracket \Gamma \vdash e:\tau_2 \rrbracket\} \\ \llbracket \Gamma \vdash \mathbf{case } e_1 \mathbf{ of } e_2 \mid e_3:\tau \rrbracket &= \{(\rho, v') \mid \rho \in \mathbf{Env}_\Gamma, (\rho, (0, v)) \in \llbracket \Gamma \vdash e_1:\tau_1 + \tau_2 \rrbracket, \\ &\quad (\rho, f) \in \llbracket \Gamma \vdash e_2:\tau_1 \rightarrow \tau \rrbracket, (v, v') \in f\} \\ &\quad \cup \{(\rho, v') \mid \rho \in \mathbf{Env}_\Gamma, (\rho, (1, v)) \in \llbracket \Gamma \vdash e_1:\tau_1 + \tau_2 \rrbracket, \\ &\quad (\rho, f) \in \llbracket \Gamma \vdash e_3:\tau_2 \rightarrow \tau \rrbracket, (v, v') \in f\} \end{aligned}$$

For recursive terms, we first notice that a recursive term $\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e$ reduces in one step to $\lambda y:\tau_1. e\{\mu x:\tau_1 \rightarrow \tau_2. e/x\}$, which is value. Therefore, the semantics of $\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e$ should be a *total* function from environments to values in $\llbracket \tau_1 \rrbracket \rightarrow \llbracket \tau_2 \rrbracket$.

This lets us derive the following equation for the semantics of $\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e$:

$$\llbracket \Gamma \vdash (\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e):\tau_1 \rightarrow \tau_2 \rrbracket \rho = \llbracket \Gamma \vdash \lambda y:\tau_1. e:\tau_1 \rightarrow \tau_2 \rrbracket \rho[x \mapsto \llbracket \Gamma \vdash (\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e):\tau_1 \rightarrow \tau_2 \rrbracket \rho]$$

which is actually a family of equations parameterized by ρ , of the form:

$$S_\rho = \llbracket \Gamma \vdash \lambda y:\tau_1. e:\tau_1 \rightarrow \tau_2 \rrbracket \rho[x \mapsto S_\rho]$$

and for each such equation a solution can be found by using the techniques we saw when studying the denotational semantics of IMP, that is, taking the functional

$$G_{\Gamma, x, y, \tau_1, \tau_2, e, \rho}(f) = \llbracket \Gamma \vdash \lambda y:\tau_1. e:\tau_1 \rightarrow \tau_2 \rrbracket \rho[x \mapsto f]$$

taking partial functions to partial functions, showing that it is continuous (it is, but it's a bear to prove directly), and therefore obtaining the solution:

$$fix(G_{\Gamma, x, y, \tau_1, \tau_2, e, \rho}) = \bigcup_{i \geq 0} G_{\Gamma, x, y, \tau_1, \tau_2, e, \rho}^i(\emptyset)$$

The final semantics of recursive terms is therefore:

$$\llbracket \Gamma \vdash (\mu x:\tau_1 \rightarrow \tau_2. \lambda y:\tau_1. e) : \tau_1 \rightarrow \tau_2 \rrbracket \rho = \text{fix}(G_{\Gamma, x, y, \tau_1, \tau_2, e, \rho})$$

As an example, work out that the semantics of the factorial function:

$$\llbracket \vdash (\mu f:\mathbf{int} \rightarrow \mathbf{int}. \lambda n:\mathbf{int}. \text{if } n = 0 \text{ then } 1 \text{ else } n \times (f \ (n - 1))) : \mathbf{int} \rightarrow \mathbf{int} \rrbracket \rho$$

yields the function given by the graph

$$\{(0, 1), (1, 1), (2, 2), (3, 6), (4, 24), \dots\}$$

5 Type inference

In the simply typed lambda calculus, we must explicitly state the type of function arguments: $\lambda x:\tau. e$. This explicitness makes it possible to type check functions.

$$\frac{\Gamma, x:\tau \vdash e:\tau'}{\Gamma \vdash \lambda x:\tau. e:\tau \rightarrow \tau'}$$

Suppose we didn't want to provide type annotations for function arguments. Consider the typing rule for functions without type annotations.

$$\frac{\Gamma, x:\tau \vdash e:\tau'}{\Gamma \vdash \lambda x. e:\tau \rightarrow \tau'}$$

The type-checking algorithm would need to guess or somehow know what type τ to put into the type context.

Can we still type check our program without these type annotations? For the simply typed lambda calculus (and many of the extensions we have considered so far), the answer is yes: we can *infer* or *reconstruct* the types of a program.

Let's consider an example to see how this type inference could work.

$$\lambda a. \lambda b. \lambda c. \text{if } a \ (b + 1) \text{ then } b \text{ else } c$$

Since the variable b is used in an addition, the type of b must be **int**. The variable a must be some kind of function, since it is applied to the expression $b + 1$. Since a has a function type, the type of the expression $b + 1$ (i.e., **int**) must be a 's argument type. Moreover, the result of the function application ($a \ (b + 1)$) is used as the test of a conditional, so it better be the case that the result type of a is also **bool**. So the type of a should be **int** \rightarrow **bool**. Both branches of a conditional should return values of the same type, so the type of c must be the same as the type of b , namely **int**.

We can write the expression with the reconstructed types:

$$\lambda a:\mathbf{int} \rightarrow \mathbf{bool}. \lambda b:\mathbf{int}. \lambda c:\mathbf{int}. \text{if } a \ (b + 1) \text{ then } b \text{ else } c$$

5.1 Constraint-based typing

We now present an algorithm that, when given a typing context Γ and an expression e , produces a set of *constraints*—equations between types (including type variables)—that must be satisfied in order for e to be well-typed in Γ .

We first introduce *type variables*, which are just placeholders for types. We use X and Y to range over type variables.

The language we will consider is the lambda calculus with integer constants and addition. We assume that all function definitions contain a type annotation for the argument, but this type may simply be a type variable X .

$$\begin{aligned} e &::= x \mid \lambda x:\tau. e \mid e_1 \ e_2 \mid n \mid e_1 + e_2 \\ \tau &::= \mathbf{int} \mid X \mid \tau_1 \rightarrow \tau_2 \end{aligned}$$

To formally define type inference, we introduce a new typing relation:

$$\Gamma \vdash e : \tau \triangleright C$$

Intuitively, if $\Gamma \vdash e : \tau \triangleright C$, then expression e has type τ provided that every constraint in the set C is satisfied. Constraints are of the form $\tau_1 = \tau_2$, intuitively meaning that types τ_1 and τ_2 are equal.

We define the judgment $\Gamma \vdash e : \tau \triangleright C$ with inference rules and axioms. When read from bottom to top, these inference rules provide a procedure that, given Γ and e , calculates τ and C such that $\Gamma \vdash e : \tau \triangleright C$.

$$\text{CT-VAR} \frac{}{\Gamma \vdash x : \tau \triangleright \emptyset} x : \tau \in \Gamma$$

$$\text{CT-INT} \frac{}{\Gamma \vdash n : \mathbf{int} \triangleright \emptyset}$$

$$\text{CT-ADD} \frac{\Gamma \vdash e_1 : \tau_1 \triangleright C_1 \quad \Gamma \vdash e_2 : \tau_2 \triangleright C_2}{\Gamma \vdash e_1 + e_2 : \mathbf{int} \triangleright C_1 \cup C_2 \cup \{\tau_1 = \mathbf{int}, \tau_2 = \mathbf{int}\}}$$

$$\text{CT-ABS} \frac{\Gamma, x : \tau_1 \vdash e : \tau_2 \triangleright C}{\Gamma \vdash \lambda x : \tau_1. e : \tau_1 \rightarrow \tau_2 \triangleright C}$$

$$\text{CT-APP} \frac{\Gamma \vdash e_1 : \tau_1 \triangleright C_1 \quad \Gamma \vdash e_2 : \tau_2 \triangleright C_2}{\Gamma \vdash e_1 e_2 : X \triangleright C_1 \cup C_2 \cup \{\tau_1 = \tau_2 \rightarrow X\}} \quad X \text{ is fresh}$$

Note that we must be careful with the choice of fresh type variables. We have omitted some of the technical details that ensure the fresh type variables in the rule CT-APP are chosen appropriately.

Let's see an example of using this typing judgment to produce a set of constraints. Consider the program $\lambda a : X. \lambda b : Y. 2 + (a (b + 3))$.

$$\frac{\frac{\frac{a : X, b : Y \vdash 2 : \mathbf{int} \triangleright \emptyset}{a : X, b : Y \vdash 2 + (a (b + 3)) : \mathbf{int} \triangleright \{Z = \mathbf{int}, X = \mathbf{int} \rightarrow Z, Y = \mathbf{int}, \mathbf{int} = \mathbf{int}\}}}{a : X \vdash \lambda b : Y. 2 + (a (b + 3)) : Y \rightarrow \mathbf{int} \triangleright \{Z = \mathbf{int}, X = \mathbf{int} \rightarrow Z, Y = \mathbf{int}, \mathbf{int} = \mathbf{int}\}}}{\vdash \lambda a : X. \lambda b : Y. 2 + (a (b + 3)) : X \rightarrow Y \rightarrow \mathbf{int} \triangleright \{Z = \mathbf{int}, X = \mathbf{int} \rightarrow Z, Y = \mathbf{int}, \mathbf{int} = \mathbf{int}\}}$$

The typing derivation means that expression $\lambda a : X. \lambda b : Y. 2 + (a (b + 3)) : X \rightarrow (Y \rightarrow \mathbf{int})$ has type $X \rightarrow Y \rightarrow \mathbf{int}$ provided that we can satisfy the constraints $Z = \mathbf{int}$, $X = \mathbf{int} \rightarrow Z$, $Y = \mathbf{int}$, and $\mathbf{int} = \mathbf{int}$.

5.2 Unification

So what does it mean for a set of constraints to be satisfied? To answer this question, we define *type substitutions* (or just *substitutions*, when it's clear from context).

5.2.1 Type substitution

A type substitution is a finite map from type variables to types. For example, we write $[X \mapsto \mathbf{int}, Y \mapsto \mathbf{int} \rightarrow \mathbf{int}]$ for the substitution that maps type variable X to \mathbf{int} , and type variable Y to $\mathbf{int} \rightarrow \mathbf{int}$.

Note that the same variable could occur in both the domain and range of a substitution. In that case, the intention is that all substitutions are performed simultaneously. For example the substitution $[X \mapsto \mathbf{int}, Y \mapsto \mathbf{int} \rightarrow X]$ maps Y to $\mathbf{int} \rightarrow X$ (not to $\mathbf{int} \rightarrow \mathbf{int}$).

More formally, we define substitution of type variables as follows.

$$\begin{aligned}\sigma(X) &= \begin{cases} \sigma(\tau) & \text{if } X \mapsto \tau \in \sigma \\ X & \text{if } X \text{ not in the domain of } \sigma \end{cases} \\ \sigma(\mathbf{int}) &= \mathbf{int} \\ \sigma(\tau \rightarrow \tau') &= \sigma(\tau) \rightarrow \sigma(\tau')\end{aligned}$$

(This is only well defined if σ is acyclic, that is, does not require the substitution of variable X within the result of substituting variable X .)

Note that we don't need to worry about avoiding variable capture, since there are no constructs in the language that bind type variables. (We'll soon see more sophisticated typing constructs that introduce binders for type variables.)

We can extend substitution to constraints, and sets of constraints in the obvious way:

$$\begin{aligned}\sigma(\tau_1 = \tau_2) &= \sigma(\tau_1) = \sigma(\tau_2) \\ \sigma(C) &= \{\sigma(c) \mid c \in C\}\end{aligned}$$

Given two substitutions σ and σ' , we write $\sigma \circ \sigma'$ for the composition of the substitutions: $\sigma \circ \sigma'(\tau) = \sigma(\sigma'(\tau))$.

5.2.2 Unification

Constraints are of the form $\tau = \tau'$. We say that a substitution σ *unifies* constraint $\tau = \tau'$ if $\sigma(\tau)$ is the same as $\sigma(\tau')$. We say that substitution σ *satisfies* (or *unifies*) set of constraints C if σ unifies every constraint in C .

For example, the substitution $\sigma = [X \mapsto \mathbf{int}, Y \mapsto \mathbf{int} \rightarrow \mathbf{int}]$ unifies the constraint $X \rightarrow (X \rightarrow \mathbf{int}) = \mathbf{int} \rightarrow Y$, since

$$\sigma(X \rightarrow (X \rightarrow \mathbf{int})) = \mathbf{int} \rightarrow (\mathbf{int} \rightarrow \mathbf{int}) = \sigma(\mathbf{int} \rightarrow Y)$$

So to solve a set of constraints C , we need to find a substitution that unifies C . More specifically, suppose that $\Gamma \vdash e : \tau \triangleright C$. Expression e is typable if and only if there is a substitution σ that satisfies C , and moreover, the type of e is $\sigma(\tau)$. If there are no substitutions that satisfy C , then we know that e is not typable.

5.2.3 Unification algorithm

To calculate solutions to constraint sets, we use the idea, due to Hindley and Milner, of using *unification* to check that the set of solutions is non-empty, and to find a “best” solution (from which all other solutions can be easily generated).

The algorithm for unification is defined as follows.

$$\begin{aligned}\text{unify}(\emptyset) &= [] \quad (\text{the empty substitution}) \\ \text{unify}(\{\tau = \tau'\} \cup C') &= \text{if } \tau = \tau' \text{ then} \\ &\quad \text{unify}(C') \\ &\quad \text{else if } \tau = X \text{ and } X \text{ not a free variable of } \tau' \text{ then} \\ &\quad \quad \text{unify}([X \mapsto \tau'](C')) \circ [X \mapsto \tau'] \\ &\quad \text{else if } \tau' = X \text{ and } X \text{ not a free variable of } \tau \text{ then} \\ &\quad \quad \text{unify}([X \mapsto \tau](C')) \circ [X \mapsto \tau] \\ &\quad \text{else if } \tau = \tau_o \rightarrow \tau_1 \text{ and } \tau' = \tau'_o \rightarrow \tau'_1 \text{ then} \\ &\quad \quad \text{unify}(C' \cup \{\tau_o = \tau'_o, \tau_1 = \tau'_1\}) \\ &\quad \text{else} \\ &\quad \quad \text{fail}\end{aligned}$$

The check that X is not a free variable of the other type ensures that the algorithm doesn't produce a cyclic substitution (e.g., $X \mapsto (X \rightarrow X)$), which doesn't make sense with the finite types that we currently have.

The unification algorithm always terminates. (Hint: what decreases on every recursive invocation? this is a bit tricky because one of the cases *adds* constraints to the set of constraints!) Moreover, it produces a solution if and only if a solution exists. The solution found is the most general solution, in the sense that if $\sigma = \text{unify}(C)$ and σ' is a solution to C , then there is some σ'' such that $\sigma' = \sigma'' \circ \sigma$.